

Example: Tree for a Query

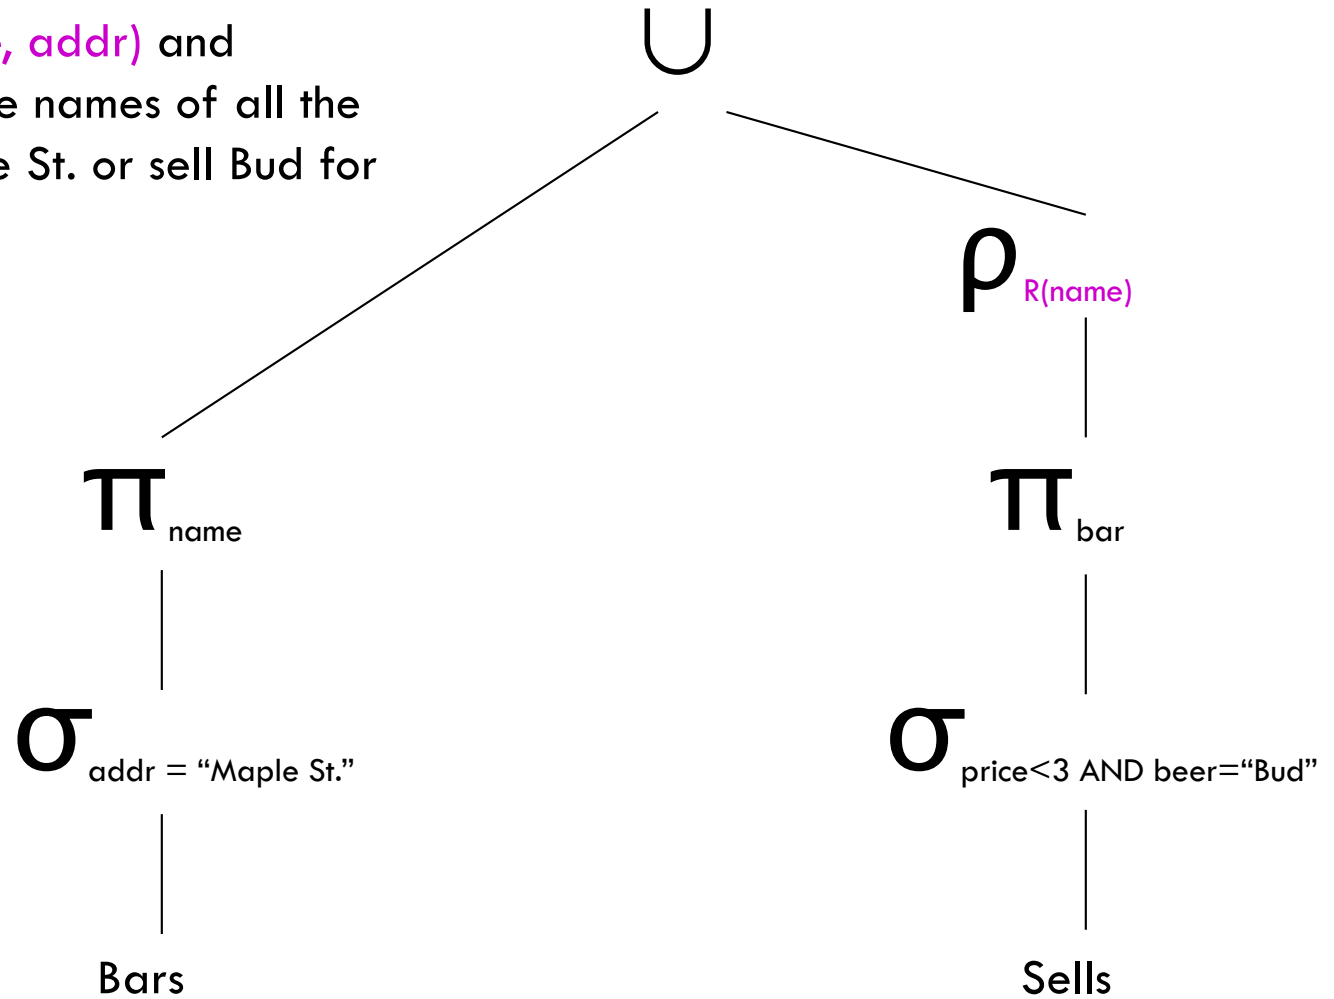
33

- Using the relations **Bars(name, addr)** and **Sells(bar, beer, price)**, find the names of all the bars that are either on Maple St. or sell Bud for less than \$3.

As a Tree:

34

Using the relations **Bars(name, addr)** and **Sells(bar, beer, price)**, find the names of all the bars that are either on Maple St. or sell Bud for less than \$3.




Example: Self-Join

35

- Using $\text{Sells}(\text{bar}, \text{beer}, \text{price})$, find the bars that sell two different beers at the same price.
- **Strategy:** by renaming, define a copy of Sells, called $\text{S}(\text{bar}, \text{beer1}, \text{price})$. The natural join of Sells and S consists of quadruples $(\text{bar}, \text{beer}, \text{beer1}, \text{price})$ such that the bar sells both beers at this price.

The Tree

36



Using $\text{Sells}(\text{bar}, \text{beer}, \text{price})$, find the bars that sell two different beers at the same price.

Schemas for Results

37

- **Union, intersection, and difference:** the schemas of the two operands must be the same, so use that schema for the result.
- **Selection:** schema of the result is the same as the schema of the operand.
- **Projection:** list of attributes tells us the schema.

Schemas for Results

38

- **Product**: schema is the attributes of both relations.
 - ▣ Distinguish two attributes with the same name.
- **Theta-join**: same as product.
- **Natural join**: union of the attributes of the two relations. Keep only one copy of the equated attributes.
- **Renaming**: the operator tells the schema.

Quiz 1

Assume set semantics

39

R

A	B
1	1
2	1
3	3

S

C	D
1	2
3	4
3	5

T

B	C
1	1
2	1
3	3

$S \bowtie T$

B	C	D
1	1	2
2	1	2
3	3	4
3	3	5

$\pi_B(R) \cap \rho_{T(B)}(\pi_C(S))$:

B
1
3

$R \bowtie (S \bowtie \rho_{T(B,C)}(R))$

A	B	C	D
1	1	1	2
2	1	1	2
3	3	3	4
3	3	3	5

The Extended Algebra

40

δ = eliminate duplicates from bags.

τ = sort tuples.

γ = grouping and aggregation.

Outerjoin : avoids “dangling tuples” = tuples that do not join with anything.

Duplicate Elimination

41

- $R1 := \delta(R2).$
- $R1$ consists of one copy of each tuple that appears in $R2$ one or more times.

Example: Duplicate Elimination

42

$R = ($

A	B
1	2
3	4
1	2

$\delta_{(R)} =$

A	B
1	2
3	4

Sorting

43

- $R1 := \tau_L(R2).$
 - ▣ L is a list of some of the attributes of $R2$.
- $R1$ is the list of tuples of $R2$ sorted first on the value of the first attribute on L , then on the second attribute of L , and so on.
 - ▣ Break ties arbitrarily.

Example: Sorting

44

$R =$ (

A	B
1	2
3	4
5	2

)

$T_B(R) =$

(

A	B
5	2
1	2
3	4

)

Aggregation Operators

45

- Aggregation operators are not formally operators of relational algebra.
- Rather, they apply to entire columns of a table and produce a single result.
- The most important examples: SUM, AVG, COUNT, MIN, and MAX.

Example: Aggregation

46

R = (

A	B
1	3
3	4
3	2

)

SUM(A) = 7

COUNT(A) = 3

MAX(B) = 4

AVG(B) = 3

Grouping Operator

47

- $R1 := \gamma_L (R2)$. L is a list of elements that are either:
 1. Individual (*grouping*) attributes.
 2. $AGG(A)$, where AGG is one of the aggregation operators and A is an attribute.
 - An arrow and a new attribute name renames the component.

Applying $\gamma_L(R)$

48

- Group R according to all the grouping attributes on list L .
 - ▣ That is: form one group for each distinct list of values for those attributes in R .
- Within each group, compute $AGG(A)$ for each aggregation on list L .
- Result has one tuple for each group:
 1. The grouping attributes and
 2. The group's aggregations.

Example: Grouping/Aggregation

49

R =

A	B	C
1	2	3
4	5	6
1	2	5

Then, average C
within groups:

A	B	X
1	2	4
4	5	6

$\gamma_{A,B,AVG(C) \rightarrow X}(R) = ??$

First, group R by A and B :

A	B	C
1	2	3
1	2	5
4	5	6

Recall: Outerjoin

50

- Suppose we join $R \bowtie_C S$.
- A tuple of R that has no tuple of S with which it joins is said to be *dangling*.
 - Similarly for a tuple of S .
- Outerjoin preserves dangling tuples by padding them NULL.

Example: Outerjoin

51

R =

A	B
1	2
4	5

S =

B	C
2	3
6	7

(1,2) joins with (2,3), but the other two tuples are dangling.

R FULL OUTERJOIN S =

A	B	C
1	2	3
4	5	NULL
NULL	6	7

Outer Join – Example

52

■ instructor

<i>ID</i>	<i>name</i>	<i>dept_name</i>
10101	Srinivasan	Comp. Sci.
12121	Wu	Finance
15151	Mozart	Music

teaches

<i>ID</i>	<i>course_id</i>
10101	CS-101
12121	FIN-201
76766	BIO-101

instructor ⋈ *teaches*

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>course_id</i>
10101	Srinivasan	Comp. Sci.	CS-101
12121	Wu	Finance	FIN-201

■ Left Outer Join

instructor ⋈_{LEFT} *teaches*

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>course_id</i>
10101	Srinivasan	Comp. Sci.	CS-101
12121	Wu	Finance	FIN-201
15151	Mozart	Music	NULL

Outer Join – Example

53

■ instructor

<i>ID</i>	<i>name</i>	<i>dept_name</i>
10101	Srinivasan	Comp. Sci.
12121	Wu	Finance
15151	Mozart	Music

teaches

<i>ID</i>	<i>course_id</i>
10101	CS-101
12121	FIN-201
76766	BIO-101

■ instructor  teaches
RIGHT

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>course_id</i>
10101	Srinivasan	Comp. Sci.	CS-101
12121	Wu	Finance	FIN-201
76766	null	null	BIO-101

■ instructor  teaches
FULL

<i>ID</i>	<i>name</i>	<i>dept_name</i>	<i>course_id</i>
10101	Srinivasan	Comp. Sci.	CS-101
12121	Wu	Finance	FIN-201
15151	Mozart	Music	null
76766	null	null	BIO-101

Operations on Bags

54

A **bag** = a set with repeated elements

All operations need to be defined carefully on bags

□ $\{a,b,b,c\} \cup \{a,b,b,b,e,f,f\} = \{a,a,b,b,b,b,c,e,f,f\}$

□ $\sigma_C(R)$: preserve the number of occurrences

□ $\Pi_A(R)$: no duplicate elimination

□ Cartesian product, join: no duplicate elimination

Important! Relational Engines work on bags, not sets!

Why Bags?

55

- SQL, the most important query language for relational databases, is actually a bag language.
- Some operations, like projection, are more efficient on bags than sets.

Operations on Bags

56

- **Selection** applies to each tuple, so its effect on bags is like its effect on sets.
- **Projection** also applies to each tuple, we do not eliminate duplicates.
- **Products** and **joins** are done on each pair of tuples, so duplicates in bags have no effect on how we operate.

Example: Bag Selection

57

R

A	B
1	2
5	6
1	2

$$\sigma_{A+B < 5} (R) =$$

A	B
1	2
1	2

Example: Bag Projection

58

R

A	B
1	2
5	6
1	2

$\pi_A(R) =$

A
1
5
1

Example: Bag Product

59

R

A	B
1	2
5	6
1	2

S

B	C
3	4
7	8

R X S =

A	R.B	S.B	C
1	2	3	4
1	2	7	8
5	6	3	4
5	6	7	8
1	2	3	4
1	2	7	8

Example: Bag Theta-Join

60

R(

A,	B)
1	2
5	6
1	2

S(

B,	C)
3	4
7	8

R $\bowtie_{R.B < S.B}$ S =

A	R.B	S.B	C
1	2	3	4
1	2	7	8
5	6	7	8
1	2	3	4
1	2	7	8

Bag Union

61

- An element appears in the union of two bags the sum of the number of times it appears in each bag.
- **Example:** $\{1,2,1\} \cup \{1,1,2,3,1\} = \{1,1,1,1,1,2,2,3\}$

Bag Intersection

62

- An element appears in the intersection of two bags the minimum of the number of times it appears in either bag
- **Example:** $\{1,2,1,1\} \cap \{1,2,1,3\} = \{1,1,2\}$.

Bag Difference

63

- An element appears in the difference $A - B$ of bags as many times as it appears in A , minus the number of times it appears in B .
- **Example:** $\{1,2,1,1\} - \{1,2,3\} = \{1,1\}$.

Beware: Bag Laws \neq Set Laws

64

- Some, but *not all* algebraic laws that hold for sets also hold for bags.
- **Example:** the commutative law for union $(R \cup S = S \cup R)$ does hold for bags.
 - ▣ Since addition is commutative, adding the number of times x appears in R and S doesn't depend on the order of R and S .

Example: A Law That Fails

65

- Set union is *idempotent*, meaning that $S \cup S = S$.
- However, for bags, if x appears n times in S , then it appears $2n$ times in $S \cup S$.
- Thus $S \cup S \neq S$ in general.
 - e.g., $\{1\} \cup \{1\} = \{1,1\} \neq \{1\}$.

What about Intersection?

Intersection is idempotent for sets and bags.