

Convolutional Neural Networks II

Swati Mishra

Applications of Machine Learning (4AL3)

Fall 2024



ENGINEERING

Review


- Convolution Operation
- Convolutional Neural Networks
- Fully Connected, ReLU, Pooling Layers
- CNN Architecture

Review

Typically, neural networks have 4 main types of layers:

- Convolutional Layer
- Pooling Layer
- Fully-Connected Layer
- ReLU Layer – element wise


Layer stacking order: CONV – RELU – POOL – FC



INPUT - CONV – RELU – POOL - CONV – RELU – POOL - ----- CONV – RELU – POOL - FC


Convolutional Neural Networks


- Convolutional layer architecture:
 - The layer comprises of filters of small width and height than the full image.
 - The number of filters extends along the depth.



Convolutional Neural Networks: Intuition

- Convolutional layer architecture :
 - The layer comprises of filters of small width and height than the full image.
 - The number of filters extends along the depth.
 - The filters are stacked behind each other, look at different features of the same **region**.




 = filter

-1	1	-1
1	0	0
-1	-1	0

This is 1 computational unit! (or 1 neuron)

Convolutional Neural Networks: Intuition


- Convolutional layer architecture :
 - The layer comprises of filters of small width and height than the full image.
 - The number of filters extends along the depth.
 - The filters are stacked behind each other, look at different features of the same **region**.
 - For forward pass, convolve the filter with the input volume (compute dot product)



$$A = aw + bx + ey + fz$$

Convolutional Neural Networks: Hyperparameters

- **Depth (D):** Number of filters used in the convolutional layer .



Convolutional Neural Networks: Hyperparameters


- **Depth (D):** Number of filters used in the convolutional layer.
- **Stride (S):** The number of pixels by which we slide the filter for convolution.

a	b	c	d
e	f	g	h
i	j	k	l

*


-1	1
0	0

Stride =1




Convolutional Neural Networks: Hyperparameters

- **Depth (D):** Number of filters used in the convolutional layer .
- **Stride (S):** The number of pixels by which we slide the filter for convolution.
- **Zero Padding (P):** Size of the zero-padding added on all 4 sides, popularly used to keep the input and output width and height are same.



Convolutional Neural Networks: Hyperparameters

- **Depth (D):** Number of filters used in the convolutional layer .
- **Stride (S):** The number of pixels by which we slide the filter for convolution.
- **Zero Padding (P):** Size of the zero-padding added on all 4 sides, popularly used to keep the input and output width and height are same.
- **Filter Size(F):** The spatial extent of each filter looking at a specific region (also called Receptive Field)




Convolutional Neural Networks: Hyperparameters

- **Depth (D):** Number of filters used in the convolutional layer .
- **Stride (S):** The number of pixels by which we slide the filter for convolution.
- **Zero Padding (P):** Size of the zero-padding added on all 4 sides, popularly used to keep the input and output width and height are same.
- **Filter Size(F):** The spatial extent of each filter looking at a specific region.


Number of computational units that fit into a layer = $(I - F + 2P)/S + 1$ where I is the input volume

Is 28 28 5


Convolutional Neural Networks: Hyperparameters



$I = 5, F = 3, S = 1, P = 0$



$I = 5, F = 3, S = 1, P = 1,$



$I = 5, F = 3, S = 2, P = 1,$

Convolutional Neural Networks: Hyperparameters

- **Depth (D):** Number of filters used in the convolutional layer .
- **Stride (S):** The number of pixels by which we slide the filter for convolution.
- **Zero Padding (P):** Size of the zero-padding added on all 4 sides, popularly used to keep the input and output width and height are same.
- **Filter Size(F):** The spatial extent of each filter looking at a specific region.



Number of computational units that fit into a layer = $(I - F + 2P)/S + 1$ where I is the input volume

4 · 5

Some configurations are impossible! For $I=10, P=0, F=3$, number of computational units not feasible.

Convolutional Neural Networks

- Three important properties of Convolutional Neural Network.:
$$s(t) = (x * w)(t)$$
 - Sparse Interactions:
 - This means every output unit does not interact with every input unit.
 - Focus on small, meaningful features such as edges with filters smaller than full image.



Convolutional Neural Networks

- Three important properties of Convolutional Neural Network.:

$$s(t) = (x * w)(t)$$


- Sparse Interactions:

- This means every output unit does not interact with every input unit.
- Focus on small, meaningful features such as edges with filters smaller than full image.
- Reduce memory requirements.
- Require fewer computations.

Convolutional Neural Networks

- Three important properties of Convolutional Neural Network.:
 - Parameter Sharing:
 - In traditional NN, each element of the weight matrix is used exactly once, multiplied by 1 element and never visited.

$$s(t) = (x * w)(t)$$





Convolutional Neural Networks

- Three important properties of Convolutional Neural Network.:

$$s(t) = (x * w)(t)$$

- Parameter Sharing:

- In traditional NN, each element of the weight matrix is used exactly once, multiplied by 1 element and never visited.
- In CNNs, we can use the same parameter for more than one function in a model.



In other words, network has **tied weights**!


Convolutional Neural Networks

- Three important properties of Convolutional Neural Network.:

$$s(t) = (x * w)(t)$$

- Parameter Sharing:

- In traditional NN, each element of the weight matrix is used exactly once, multiplied by 1 element and never visited.
- In CNNs, we can use the same parameter for more than one function in a model.



F=11
S=4
P=0
D=96

Output of this layer =

$$(I - F + 2P)/S + 1 =$$

55x55x96

All 96 units are connected to a region of 11x11x3


Convolutional Neural Networks

- Three important properties of Convolutional Neural Network.:

$$s(t) = (x * w)(t)$$

- Parameter Sharing:

- In traditional NN, each element of the weight matrix is used exactly once, multiplied by 1 element and never visited.
- In CNNs, we can use the same parameter for more than one function in a model.



F=11
S=4
P=0
D=96

Output of this layer =

$$(I - F + 2P)/S + 1 =$$

55x55x96

Each computational unit has 11x11x3 weights and 1 bias

All 96 units are connected to a region of 11x11x3


Convolutional Neural Networks

- Three important properties of Convolutional Neural Network.:

$$s(t) = (x * w)(t)$$

- Parameter Sharing:

- In traditional NN, each element of the weight matrix is used exactly once, multiplied by 1 element and never visited.
- In CNNs, we can use the same parameter for more than one function in a model.



F=11
S=4
P=0
D=96

Output of this layer =

$$(I - F + 2P)/S + 1 =$$

55x55x96 ~~1.1~~ ✓

Total parameters in this layer?




Convolutional Neural Networks

- Three important properties of Convolutional Neural Network.:

$$s(t) = (x * w)(t)$$

- Parameter Sharing:

- In traditional NN, each element of the weight matrix is used exactly once, multiplied by 1 element and never visited.
- In CNNs, we can use the same parameter for more than one function in a model.



We assume:
If one feature is useful to compute at some spatial position, then it should also be useful to compute at a different but related position as well.


Convolutional Neural Networks

- Three important properties of Convolutional Neural Network.:

$$s(t) = (x * w)(t)$$

- Parameter Sharing:

- In traditional NN, each element of the weight matrix is used exactly once, multiplied by 1 element and never visited.
- In CNNs, we can use the same parameter for more than one function in a model.



We constrain the computational units in each depth slice to use the same weights and bias.


Convolutional Neural Networks

- Three important properties of Convolutional Neural Network.:

$$s(t) = (x * w)(t)$$

- Parameter Sharing:

- In traditional NN, each element of the weight matrix is used exactly once, multiplied by 1 element and never visited.
- In CNNs, we can use the same parameter for more than one function in a model.



$$F=11, S=4, P=0, D=96$$

Output of this layer = 55x55x96

Without parameter sharing: $55 \times 55 \times 96 \times (11 \times 11 \times 3 + 1)$

With parameter sharing: $96 \times (11 \times 11 \times 3 + 1)$

All 96 units are connected to a region of 11x11x3

Convolutional Neural Networks

- Three important properties of Convolutional Neural Network.:

$$s(t) = (x * w)(t)$$

- Parameter Sharing:

- In traditional NN, each element of the weight matrix is used exactly once, multiplied by 1 element and never visited.
 - In CNNs, we can use the same parameter for more than one function in a model.
 - In practice during backpropagation, every unit computes the gradient for its weights,
 - However, these gradients will be added up across each depth slice
 - The gradients only update a single set of weights per slice.
 - Only because of this technique, can we apply convolution during forward pass.

Convolutional Neural Networks

- Three important properties of Convolutional Neural Network.:

$$s(t) = (x * w)(t)$$

- Parameter Sharing:

- In traditional NN, each element of the weight matrix is used exactly once, multiplied by 1 element and never visited.
- In CNNs, we can use the same parameter for more than one function in a model.
- In practice during backpropagation, every unit computes the gradient for its weights,
- However, these gradients will be added up across each depth slice
- The gradients only update a single set of weights per slice.
- Only because of this technique, can we apply convolution during forward pass.



Convolutional Neural Networks

- Three important properties of Convolutional Neural Network.: $s(t) = (x * w)(t)$
 - Equivariant Representations:
 - The parameter sharing causes the property of equivariance to translation.
 - Equivariant means that if the input changes, the output changes in the same way.
 - If g translates the input (or shifts it), then convolution function is equivariant to g

$$f(g(x)) = g(f(x))$$

Convolutional Neural Networks

- Three important properties of Convolutional Neural Network.: $s(t) = (x * w)(t)$
 - Equivariant Representations:
 - The parameter sharing causes the property of equivariance to translation.
 - Equivariant means that if the input changes, the output changes in the same way.
 - If g translates the input (or shifts it), then convolution function is equivariant to g

$$f(g(x)) = g(f(x))$$

Let's assume an image $I(x, y)$

- *shifting transformation $I'(x, y) = I(x - 1, y)$*

Transformation to I + convolution

=

Convolution to I' + Transformation

Convolutional Neural Networks: Example

0	0	0	0	0	0	0
0	1	1	0	0	2	0
0	1	1	0	1	0	0
0	1	1	1	2	1	0
0	1	0	2	0	2	0
0	0	0	2	1	0	0
0	0	0	0	0	0	0

`x[0, :, :]`

0	0	0	0	0	0	0
0	0	0	2	0	1	0
0	1	1	0	1	1	0
0	2	0	2	2	0	0
0	1	1	1	2	0	0
0	0	1	2	2	0	0
0	0	0	0	0	0	0

`x[1, :, :]`

✓ 0.0s

0	0	0	0	0	0	0
0	0	0	2	0	1	0
0	1	1	0	1	1	0
0	2	0	2	2	0	0
0	1	1	1	2	0	0
0	0	1	2	2	0	0
0	0	0	0	0	0	0

`x[2, :, :]`

✓ 0.0s

```
x = torch.rand(3, 5, 5)
```

Convolutional Neural Networks: Example

0	0	0	0	0	0	0
0	1	1	0	0	2	0
0	1	1	0	1	0	0
0	1	1	1	2	1	0
0	1	0	2	0	2	0
0	0	0	2	1	0	0
0	0	0	0	0	0	0

0	0	0	0	0	0	0
0	0	0	2	0	1	0
0	1	1	0	1	1	0
0	2	0	2	2	0	0
0	1	1	1	2	0	0
0	0	1	2	2	0	0
0	0	0	0	0	0	0

0	0	0	0	0	0	0
0	0	0	2	0	1	0
0	1	1	0	1	1	0
0	2	0	2	2	0	0
0	1	1	1	2	0	0
0	0	1	2	2	0	0
0	0	0	0	0	0	0

Filter F0

1	1	1
1	0	-1
-1	0	1

`f[0,:,:]`
✓ 0.0s

-1	0	-1
0	1	0
0	0	0

`f[1,:,:]`
0.0s

1	1	-1
0	1	0
-1	0	0

`f[2,:,:]`
✓ 0.0s

0	7	3
4	4	6
8	8	5

Output

Readings

Required Readings:

Introduction to Statistical Learning

- Chapter 10 – Section 10.3 page 406 - 412

Supplemental Readings (Not required but recommended):

Deep Learning

- Chapter 9 – page 330 - 340

Thank You
