

# Operating Systems: CPU Scheduling

Neerja Mhaskar

Department of Computing and Software, McMaster University, Canada

**Acknowledgements:** Material based on the textbook Operating Systems Concepts (Chapter 5)

# Basic Concepts

- OS schedules almost all resources available to the system.
- CPU being the primary resource, scheduling processes to execute on the CPU is central to OS design.
- Maximizing CPU utilization, is feasible as processes alternate using CPU and waiting for I/O

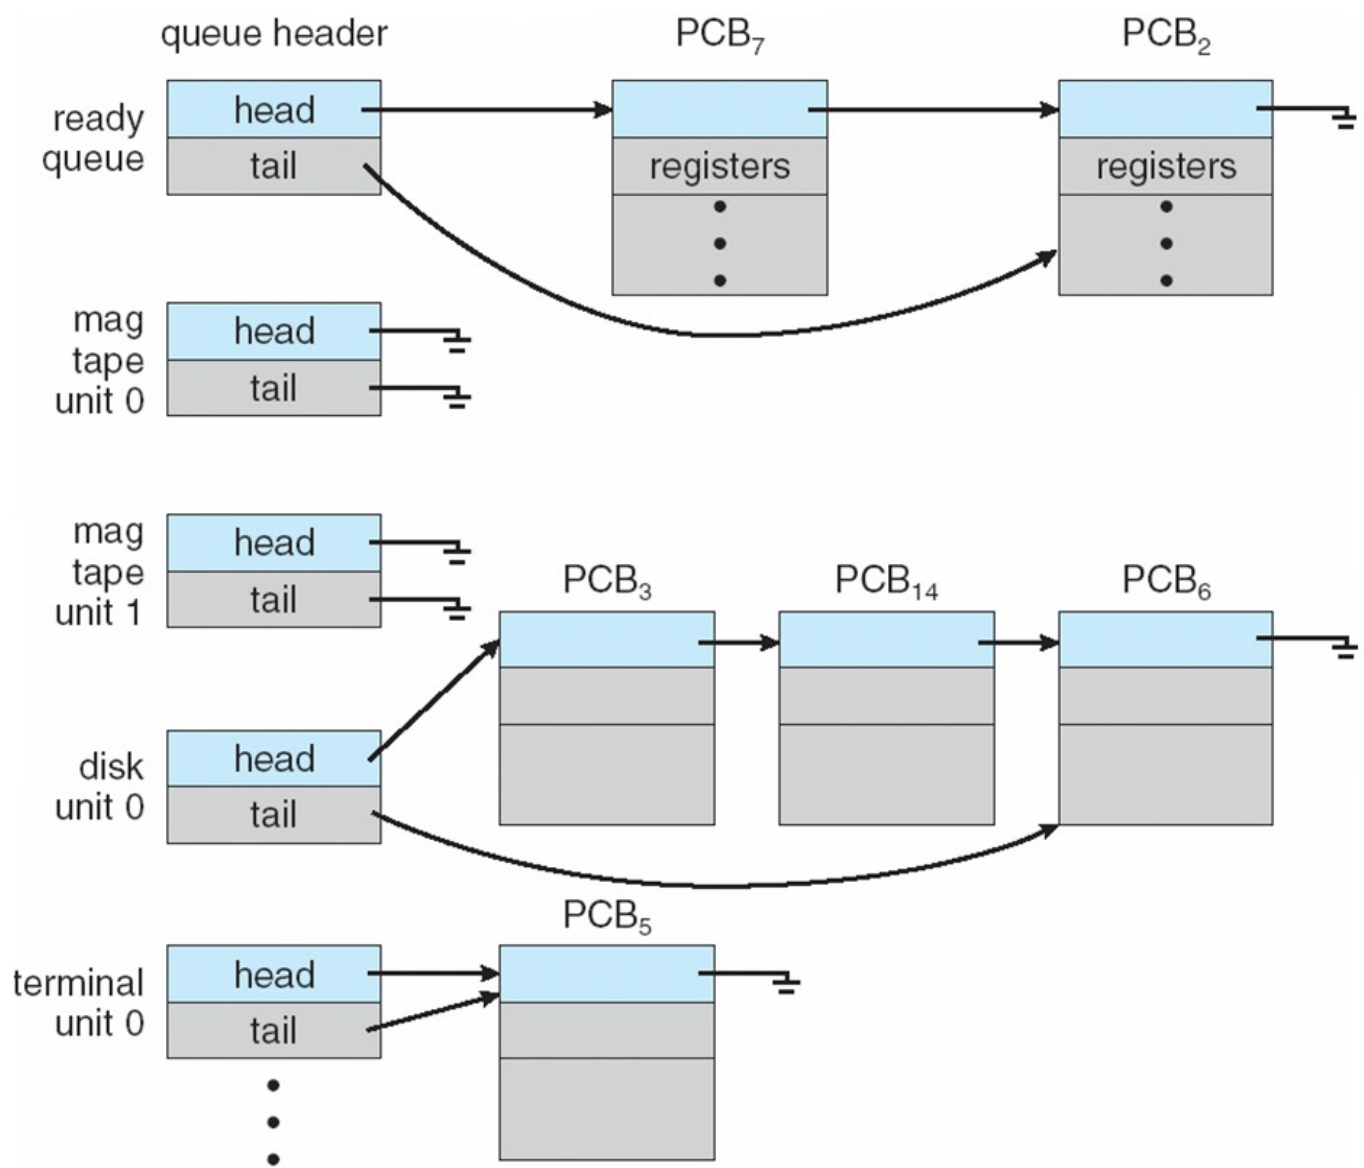
# CPU burst and I/O burst

- Process execution consists of a **cycle** of CPU execution (**CPU burst**) and I/O wait (**I/O burst**).
- An I/O-bound program typically has many short CPU bursts.
- A CPU-bound program might have a few long CPU bursts
- This distribution is important in the selection of an appropriate CPU-scheduling algorithm.

# Process Scheduling

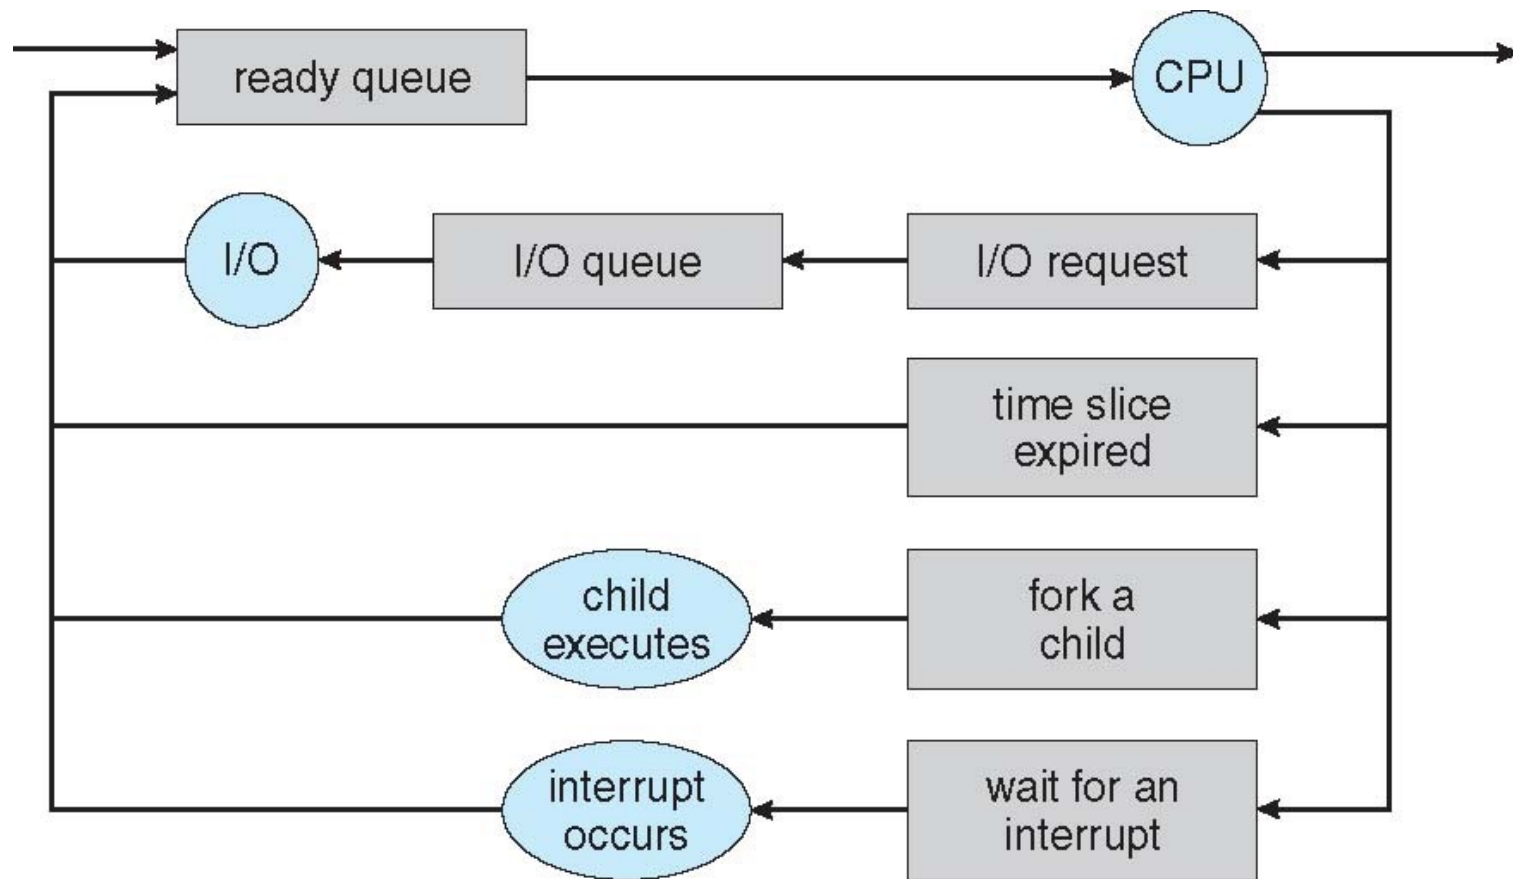
- Operating System maintains **scheduling queues** of processes
  - **Job queue** – set of all processes in the system
  - **Ready queue** – set of all processes residing in main memory, ready and waiting to execute
  - **Device queues** – set of processes waiting for an I/O device.  
Usually a separate device queue for each device.
- Processes migrate among the various queues
- **Process scheduler**: selects a process from a queue by implementing appropriate scheduling algorithm.

# Ready Queue And Various I/O Device Queues



# Representation of Process Scheduling

**Queueing diagram** represents queues, resources, flows



# CPU scheduler

- **CPU scheduler (or Short-term scheduler)** – selects from among the processes in ready queue, and allocates the CPU to one of them (see next slide)
  - Sometimes the only scheduler in a system
  - Short-term scheduler is invoked frequently (milliseconds)  $\Rightarrow$  (must be fast)

# Non-preemptive Vs. Preemptive

- **Non-preemptive scheduling** - a running process is executed till completion without interruption.
- **Preemptive scheduling** – a running process may be interrupted and moved to the Ready queue by the OS.




# Dispatcher

- Dispatcher module gives control of the CPU to the process selected by the short-term scheduler; this involves:
  - switching context
  - switching to user mode
  - jumping to the proper location in the user program to restart that program

# Dispatcher

- **Dispatch latency** – time taken by the dispatcher to stop one process and start put another process onto the CPU.



# Other Scheduler

- **Long-term scheduler** (or **job scheduler**) – selects which processes should be brought into the ready queue
  - Long-term scheduler is invoked infrequently (seconds, minutes) ⇒ (may be slow). Therefore, can use sophisticated scheduling algorithms.
  - Typically, seen on a process intensive systems.
  - The long-term scheduler controls the **degree of multiprogramming** and good ***process mix***

# Scheduling Criteria

- **CPU utilization** – keep the CPU as busy as possible
- **Throughput** – # of processes that complete their execution per time unit
- **Turnaround time** – amount of time to execute a particular process
  - Turnaround time for a process = waiting time + CPU burst time
- **Waiting time** – amount of time a process has been waiting in the ready queue

## Scheduling Algorithm Optimization Criteria

- Maximize CPU utilization and throughput
- Minimize turnaround time and waiting time

# Gantt Chart

**Gantt Chart:** is a bar chart that illustrates a particular **process schedule**, including the start and finish times of each of the participating processes.



# Gantt Chart Example

Suppose processes  $P_1$ ,  $P_2$ ,  $P_3$  have CPU burst time 24, 3, 3 respectively and arrive in the same order.

**The Gantt Chart for the schedule is:**



- **Waiting time** for  $P_1 = 0$ ;  $P_2 = 24$ ;  $P_3 = 27$  msec
- **Average waiting time:**  $(0 + 24 + 27)/3 = 17$  msec
- **Turnaround time for a process = waiting time + CPU burst time**
- **Average turnaround time:**  $(24 + 27 + 30) / 3 = 27$  msec

# Scheduling Algorithms

- First Come First Serve (FCFS) Scheduling
- Shortest Job First (SJF) Scheduling
  - Shortest Remaining Time First Scheduling
- Priority Scheduling
- Round Robin Scheduling

# First- Come, First-Served (FCFS) Scheduling

- Process requesting the CPU first is allocated the CPU first.
  - The implementation of the FCFS policy can be achieved with a FIFO queue.
- FCFS scheduling algorithm is non-preemptive
- **Disadvantage:** Average CPU waiting time for a process to use CPU is long.



# FCFS Scheduling (Cont.)

- Suppose that the processes arrive in the order:

$$P_2, P_3, P_1$$

The Gantt chart for the schedule is:



- Waiting time for  $P_1 = 6$  msec;  $P_2 = 0$  msec;  $P_3 = 3$  msec
- Average waiting time:  $(6 + 0 + 3)/3 = 3$  msec
- **Much better than previous case!**
- **Convoy effect** - short process behind long process
  - Results in lower resource utilization


# Shortest-Job-First (SJF) Scheduling

- Associated with each process is **the length of its next CPU burst**
- **SJF** use these lengths to schedule the process with the shortest CPU burst.
  - If the next CPU bursts of two processes are the same, FCFS scheduling is used to break the tie.
- **SJF is optimal** – gives minimum average waiting time for a given set of processes

# Example of SJF

<u>Process</u>	<u>CPU Burst Time</u>
$P_1$	6
$P_2$	8
$P_3$	7
$P_4$	3

## ■ SJF scheduling chart



■ Average waiting time =  $(0 + 3 + 9 + 16) / 4 = 7$  msec

■ Turnaround time =  $(3 + 9 + 16 + 24) / 4 = 13$  msec


# Shortest Remaining Time First

- SJF algorithm can either be preemptive or nonpreemptive
- Preemptive version of SJF is called **shortest-remaining-time-first**
  - At a given time, the process with shortest remaining time is scheduled first.
- Now we add the concepts of varying arrival times and preemption to the analysis

# Example of Shortest-remaining-time-first

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
$P_1$	0	8
$P_2$	1	4
$P_3$	2	9
$P_4$	3	5

## *Preemptive* SJF Gantt Chart



- Average waiting time =  $[(10-1)+(1-1)+(17-2)+(5-3)]/4 = 26/4 = 6.5$  msec

# Priority Scheduling

- A **priority number** (integer) is associated with each process
  - Equal priority process scheduled in FCFS order.
- The CPU is allocated to the process with the highest priority (**smallest integer  $\equiv$  highest priority**)
- Priority scheduling can either be preemptive or non-preemptive
- Problem  $\equiv$  **Starvation** – low priority processes may never execute
  - **Solution  $\equiv$  Aging** – as time progresses increase the priority of the process

# Example of Priority Scheduling

<u>Process</u>	<u>Burst Time(ms)</u>	<u>Priority</u>
$P_1$	10	3
$P_2$	1	1
$P_3$	2	4
$P_4$	1	5
$P_5$	5	2

## ■ Priority scheduling Gantt Chart



## ■ Average waiting time = $(0 + 1 + 6 + 16 + 18) / 5 = 8.2$ ms (milliseconds)

# Round Robin (RR)

- Similar to FCFS scheduling, but preemption is added to enable the system to switch between processes.
- Processes are scheduled on FCFS basis from the ready queue, where each process gets a small unit of CPU time (**time quantum  $q$** )
- After this time has elapsed, the process is preempted and added to the end of the ready queue.
- Timer interrupts every quantum to schedule next process
- **Preemptive Scheduling**



# Example of RR with Time Quantum = 4


<u>Process</u>	<u>Burst Time</u>
----------------	-------------------

$P_1$	24
-------	----

$P_2$	3
-------	---

$P_3$	3
-------	---

- The Gantt chart is:



- Typically, higher average *turnaround time* than SJF, but better *response time*.

# Round Robin Performance

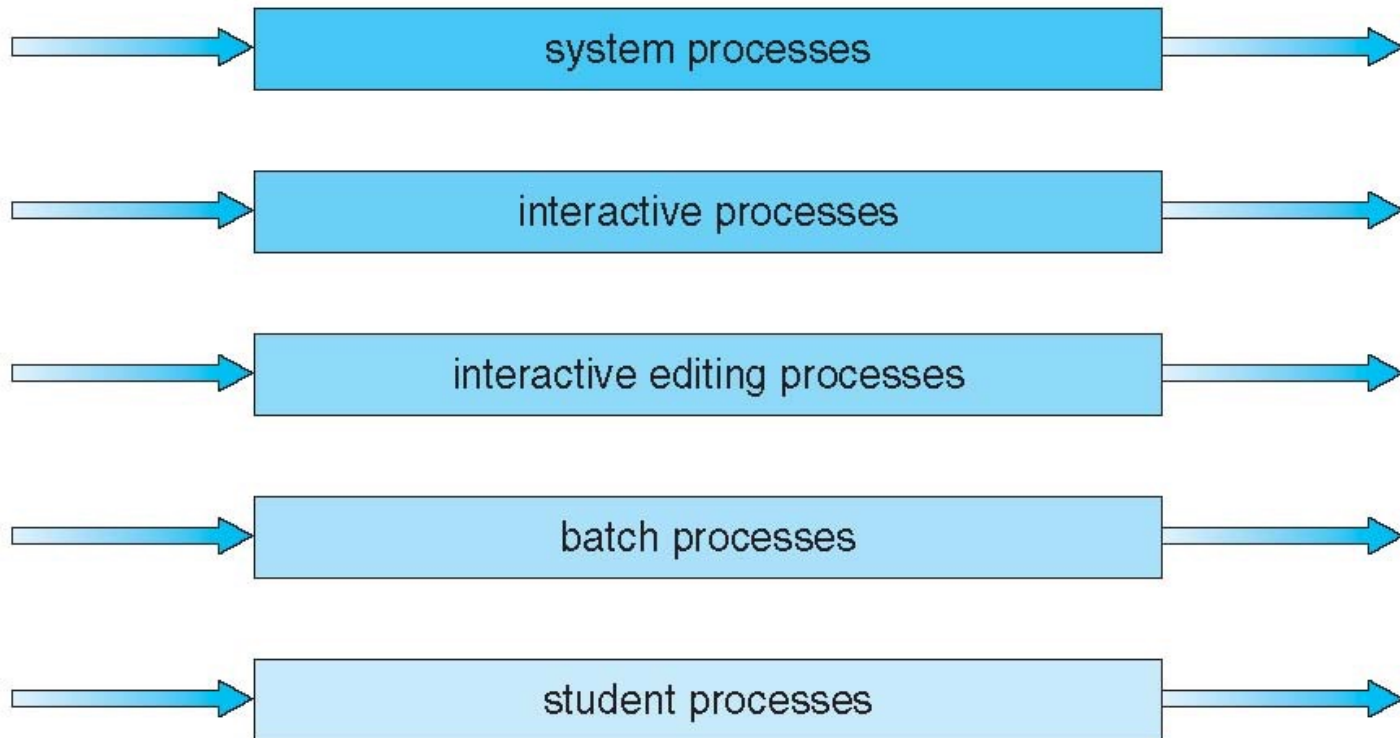
- $q$  large  $\Rightarrow$  FIFO
- $q$  must be large with respect to context switching time, otherwise overhead is too high.
  - $q$  is usually 10 to 100 milli seconds and context switch  $< 10$  micro second.

# Multilevel Queue and Multilevel Feedback Queue Scheduling

- **Multilevel Queue Scheduling:** Ready queue is partitioned into various separate queues. Process resides permanently in a given queue.
  - Each queue has its own scheduling algorithm
  - Scheduling must be done between the queues:
  - ***Fixed priority preemptive scheduling:*** Serve processes of highest priority first.
    - ***Possibility of starvation.***
  - ***Time slice*** – each queue gets a certain amount of CPU time which it can schedule amongst its processes.

# Multilevel Queue Scheduling

highest priority



lowest priority

# Multilevel Feedback Queue Scheduling

## Multilevel Feedback Queue Scheduling - Disadvantage

- **Inflexible** – as processes are permanently assigned to a given queue.
- In contrast, in **Multilevel Feedback Queue Scheduling** a process can move between the various queues.
  - Process in low priority queue can be moved to high priority queue, and aging can be implemented this way.

# Thread Scheduling

- When threads are supported by the kernel, threads scheduled not processes
- To run on CPU user level threads must be mapped to an associated kernel level thread.
- POSIX Pthreads allows setting scheduling parameters during thread creation.