

# Operating Systems: Synchronization Tools and Examples – Part II

Neerja Mhaskar

Department of Computing and Software, McMaster University, Canada

**Acknowledgements:** Material based on the textbook Operating Systems Concepts (Chapters 6 & 7)

# Semaphore

- Semaphores is a synchronization tool that provides more sophisticated ways (other than Mutex locks) for processes to synchronize their activities.
- Semaphore **S** – integer variable (usually initialized with a +ve integer)
- Can only be accessed via two indivisible (**atomic**) operations
  - **wait()** and **signal()**
    - Originally called **P()** and **V()**

# Semaphore – Wait(S) and Signal() Operations

- **Wait() operation:** Decreases the Semaphore (S) value only if  $S > 0$

```
wait(S) {  
    while (S <= 0); // busy wait  
    S--;  
}
```

- **Signal() operation:** Increases the Semaphore (S) value.

```
signal(S) {  
    S++;  
}
```

# Semaphore Usage

- Semaphores can solve various synchronization problems
  - Used as locks (*binary semaphores*)
  - Used to control access to a given resource consisting of a finite number of instances (*counting semaphore*)
  - To solve synchronization problems (may or may not access shared data)

# Binary Semaphore

- **Binary semaphore** – integer value (S) can range only between 0 and 1
  - Initially, semaphore  $S = 1$  (lock is available),  $S = 0$  (lock is unavailable).
  - Behaves like a mutex lock.

```
do {  
  
    wait(S);  
  
    /*critical section*/  
  
    Signal(S);  
  
    /*remainder section*/  
  
} while (TRUE);
```

# Counting Semaphore

- **Counting semaphore** – integer value ( $S$ ) can range over an unrestricted domain.
  - Used to control access to finite instances ( $N$ ) of a resource.
  - Initially,  $S = N$  (*all instances of resource available*)
  - To acquire the resource: processes perform `wait(S)`
  - To release the resource (after use) - processes perform `signal(S)`
  - $S = 0$ , *implies all instances of the resource are being used*
  - Waiting processes *block or is busy waiting* until  $S > 0$  (*an instance of resource is available*)

# Semaphores – synchronization problems

Question 1: Consider two concurrently running processes  $P_1$  and  $P_2$  that require  $P_1$  to execute before  $P_2$  (here  $P_1$  and  $P_2$  may or may not access/modify shared variables.). Use semaphores to synchronize processes  $P_1$  and  $P_2$ .

# Semaphores – synchronization problems

## Solution to Question 1:

- Create a semaphore.
- Initialize `synch = 0`

**P1:**

`S1;`

`signal(synch);`

**P2:**

`wait(synch);`

`S2;`



# Semaphores – synchronization problems Example

Question 2: Illustrate how you would use semaphores to synchronize processes  $P_1$ ,  $P_2$ ,  $P_3$ ,  $P_4$  where  $P_1$  is executed before  $P_2$ ,  $P_2$  is executed before  $P_3$ ,  $P_3$  is executed before  $P_4$ .

# Semaphores – synchronization problems Example

**Solution to Question 2:** Initialize  $\text{synch}_1, \text{synch}_2, \text{synch}_3 = 0$

**P<sub>1</sub>:**

```
S1;  
signal(synch1);
```

**P<sub>2</sub>:**

```
wait(synch1);  
S2;  
signal(synch2)
```

**P<sub>3</sub>:**

```
wait(synch2);  
S3;  
signal(synch3);
```

**P<sub>4</sub>:**

```
wait(synch3);  
S4;
```

Semaphores studied so far suffer  
from busy waiting!

# Semaphore Implementation with no Busy waiting

- With each semaphore there is an associated waiting queue
- Each semaphore has **two data items**:
  - **value (of type integer)**
  - **pointer to the list of processes waiting**

- **Semaphore Structure**

```
typedef struct{  
  
    int value;  
  
    struct process *list;  
  
} semaphore;
```

- `wait()` and `signal()` operations are atomic.
- **Two additional operations: `block()` and `wakeup()`**

# Semaphore Implementation with no Busy waiting

## Wait Operation:

```
wait(semaphore *S) {  
    S->value--;  
  
    if (S->value < 0) {  
        add this process to S->list;  
  
        block();  
    }  
}
```

**block()** – places the process invoking the operation on the appropriate waiting queue

## Signal Operation:

```
signal(semaphore *S) {  
    S->value++;  
  
    if (S->value <= 0) {  
        remove a process P from S->list;  
  
        wakeup(P);  
    }  
}
```

**wakeup()** – removes a processes in the waiting queue and places it in the ready queue

## Semaphore Implementation with no Busy waiting

- Note that in this implementation, semaphore values may be negative
- If a semaphore value is negative, its magnitude is the number of processes waiting on that semaphore.
- Implementation of semaphores using waiting queues may result in **Deadlock and/or starvation**.

# Deadlock and Starvation

- **Deadlock** – two or more processes are waiting indefinitely for an event that can be caused by only one of the waiting processes
- Let  $S$  and  $Q$  be two semaphores initialized to 1

$P_0$	$P_1$
<code>wait(S);</code>	<code>wait(Q);</code>
<code>wait(Q);</code>	<code>wait(S);</code>
<code>...</code>	<code>...</code>
<code>signal(S);</code>	<code>signal(Q);</code>
<code>signal(Q);</code>	<code>signal(S);</code>

- **Starvation – indefinite blocking**
  - A process may never be removed from the semaphore queue in which it is suspended

# Semaphores in C

- Semaphores are not part of Pthread standards and instead they belong to the POSIX SEM extension.
- To use semaphores in C on a Linux machine include 'semaphore.h' header file:
  - `#include <semaphore.h>`
  - `sem_t` data type for semaphores.
  - semaphores created with `sem_init()` function.
    - `sem_init()` - Takes 3 arguments:
      1. Pointer to the semaphore
      2. Flag indicating the level of sharing
      3. The semaphore's initial value
  - `sem_wait()` ;
  - `sem_post()` ;



# Semaphores in C

```
/*Declaring Semaphore*/
```

```
Sem_t sem;
```

```
/*Initialize Semaphore*/
```

```
if (sem_init(&sem, 0, n) !=0){
```

```
    printf("Error in initializing empty semaphore \n"
```

```
}
```

```
/* acquire the semaphore */
```

```
sem_wait(&sem);
```

```
/* critical section */
```

```
/* release the semaphore */
```

```
Sem_post(&sem);
```

## **sem\_init** arguments

1. pointer to the semaphore
2. flag indicating the level of sharing
3. The semaphore's initial value

Returns 0 when semaphore created with no errors, otherwise returns non zero value.

# Problems with Semaphores

- Incorrect use of semaphore operations could result in mutual exclusion being violated, deadlock and/or starvation.
- Consider the following situation:
  - Processes share a Semaphore `mutex`, which is initialized to 1.
- Suppose that a process interchanges the order in which the `wait()` and `signal()` operations on the semaphore `mutex` are executed. That is, it executes

```
signal(mutex);  
  
...  
  
critical section  
  
...  
  
wait(mutex);
```

- **What is the problem with the above code?**

# Problems with Semaphores

- Suppose that a process replaces `signal(mutex)` with `wait(mutex)`.

That is, it executes

```
wait(mutex);  
  
...  
  
critical section  
  
...  
  
wait(mutex);
```

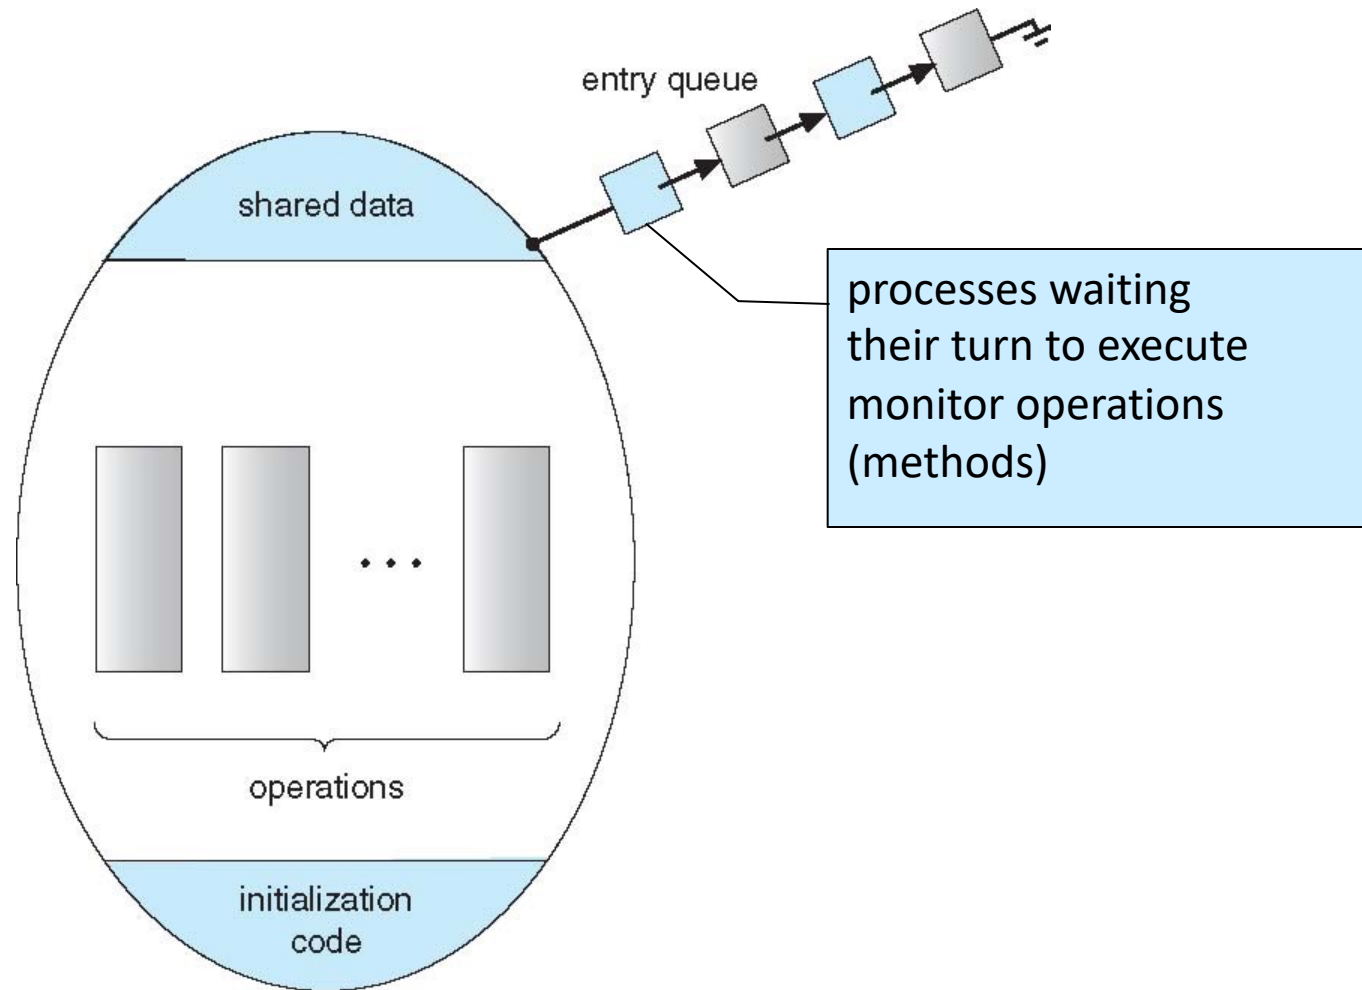
## What is the problem here?

- Omitting of `wait(mutex)` or `signal(mutex)` (or both)
  - Deadlock and starvation are possible.

# Monitors

- **Monitor** is a high-level abstraction that provides a convenient and effective mechanism for process synchronization.
- A **monitor type** is an ADT has
  - Private data members
  - Public functions that are implicitly executed with mutual exclusion.
    - For each monitor instance there exists a mutual exclusion lock.
    - To enter the monitor, a process acquires the mutual exclusion lock
    - While exiting the monitor, a process releases the lock, and therefore the monitor, for other threads.

# Schematic view of a Monitor



# Syntax of Monitor

```
monitor monitor-name  
{  
    shared variable declarations  
  
    procedure P1 (...) { ... }  
    .  
    .  
    .  
    procedure Pn (...) {.....}  
  
    initialization code (...) { ... }  
  
}
```

# Monitor Example

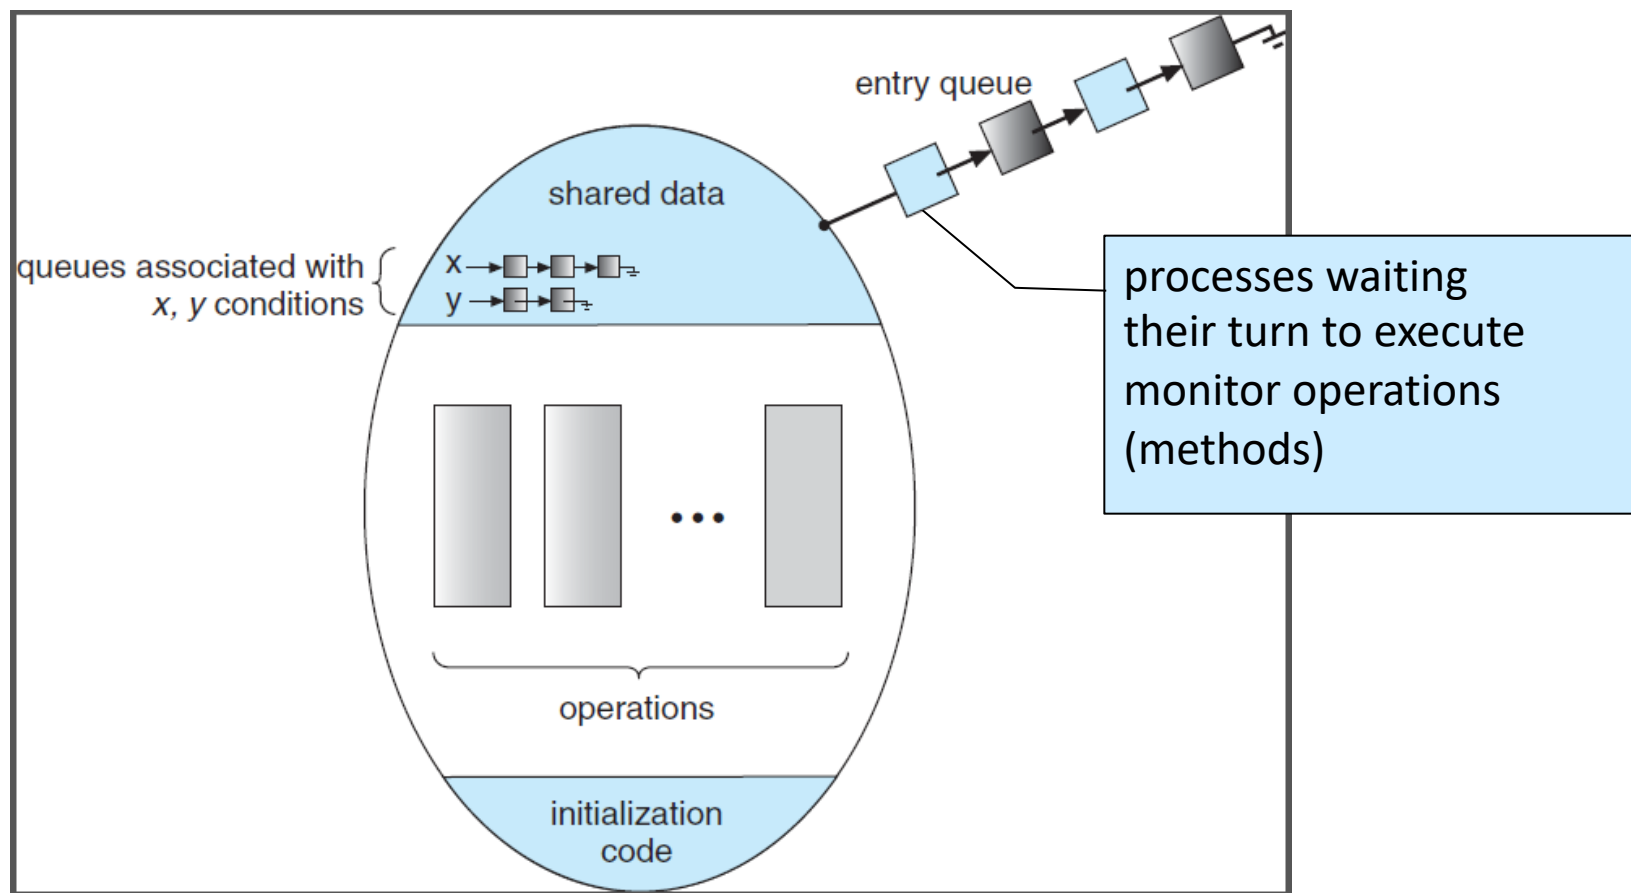
```
monitor BankAccount {  
    private: int amount;  
    void BankAccount_init(){  
        amount = 0;  
    }  
    void withdraw(int value){  
        amount = amount - value;  
    }  
    void deposit(int value){  
        amount = amount + value;  
    }  
};
```

# Condition Variables

- Monitor construct defined so far is not powerful for modelling some synchronization schemes.
- Condition variables (CV) are additional mechanisms to achieve the above and defined as: `condition x, y;`
- They have legal atomic operations: `wait()` and `signal()`
- Operations accessed as `x.wait()`, `x.signal()`
- `x.wait()` – blocks the process and adds it to a queue associated with x.
- `x.signal()` – wakes up exactly one process from x's queue of waiting processes.
  - If no processes waiting, then does nothing.



# Monitor with Condition Variables



# Monitor Example with Condition Variables

- How to ensure the following two conditions are met:
  - Deposits do not take place if amount  $\geq 500$
  - Withdrawals do not take place if amount  $\leq 0$
- Assume that the value deposited/withdrawn is the same.

```

monitor BankAccount {
    int amount;
    condition c1, c2;

    void BankAccount_init(){
        amount = 0;
    }
    void withdraw(int value){
        if(amount <=0){c1.wait();}
        amount = amount - value;
        c2.signal();
    }
    void deposit(int value){
        if(amount >=500){c2.wait();}
        amount = amount + value;
        c1.signal();
    }
};

```

# Condition Variables Choices

- If process P invokes `x.signal()`, and process Q is suspended in `x.wait()`, what should happen next?
  - Both Q and P **cannot execute in parallel in the monitor**. If Q is resumed, then P must wait
- Options include
  - **Signal and wait** – P waits until Q either leaves the monitor or it waits for another condition
  - **Signal and continue** – Q waits until P either leaves the monitor or it waits for another condition
  - Both have pros and cons – language implementer can decide

# Resuming Processes within a Monitor

- If several processes queued on condition `x`, and `x.signal()` executed, which should be resumed?
  - FCFS (first come first serve) frequently not adequate
- **conditional-wait** construct of the form `x.wait(c)`
  - where `c` is **priority number**
  - process with ***lowest number*** (highest priority) is scheduled next

# Classical Problems of Synchronization

- Bounded-Buffer Problem (Producer - Consumer Problem) (discussed in part I)
- Readers and Writers Problem
- Dining-Philosophers Problem

# Bounded-Buffer Problem – Semaphore Solution

- $n$  buffers, each can hold one item
- Semaphore **mutex** initialized to the value 1
  - `mutex = 1` => indicates that `mutex` lock is available
  - `mutex = 0` => indicates that `mutex` lock is unavailable
- Semaphore **full** initialized to the value 0, as initially the number of filled slots in the buffers is zero.
- Semaphore **empty** initialized to the value  $n$ , as initially the number of empty slots in the buffer is  $n$ .

# Bounded Buffer Problem – Semaphore Solution (Cont.)

- The structure of the producer process

```
do {  
  
    ...  
    /* produce an item in next_produced */  
  
    ...  
  
    wait(empty);  
  
    wait(mutex);  
  
    ...  
    /* add next produced to the buffer */  
  
    ...  
  
    signal(mutex);  
  
    signal(full);  
  
} while (true);
```




# Bounded Buffer Problem Semaphore Solution (Cont.)

- The structure of the consumer process

```
Do {  
  
    wait(full) ;  
  
    wait(mutex) ;  
  
    ...  
  
    /* remove an item from buffer to next_consumed */  
  
    ...  
  
    signal(mutex) ;  
  
    signal(empty) ;  
  
    ...  
  
    /* consume the item in next consumed */  
  
    ...  
  
} while (true);
```

# Dining-Philosophers Problem

- Philosophers spend their lives alternating between **thinking** and **eating**
- Don't interact with their neighbors
- Occasionally try to pick up 2 chopsticks (**one at a time**) to eat from a bowl
- **Need both chopsticks to eat**, then release both when done



# Dining-Philosophers Problem Semaphore Solution

In the case of 5 philosophers

➤ Shared data

- Bowl of rice (data set) – assumed infinite
- Semaphore **chopstick [5]** initialized to 1 (indicates its available)



# Dining-Philosophers Problem Algorithm Semaphore Solution

The structure of Philosopher  $i$ :

```
do {  
    wait (chopstick[i] );  
    wait (chopstick[ (i + 1) % 5]);  
  
        // eat  
  
    signal (chopstick[i] );  
    signal (chopstick[ (i + 1) % 5]);  
  
        // think  
} while (TRUE);
```



What is the problem with this algorithm?

# Dining-Philosophers Problem Algorithm

- Consider the following scenario in which all the philosophers are hungry and try to eat at the same time.
- Each philosopher grabs the chopstick to their right (`chopstick[i]`).
- Each semaphore `chopstick[i]=0`, where  $0 \leq i \leq 4$ .
- When each philosopher tries to grab her left chopstick (`chopstick[i+1] % 5`), she will be delayed forever.  
Thus resulting in a *deadlock*.

## Dining-Philosophers Problem Algorithm (Cont.)

- Deadlock can be handled by
  - Allowing at most 4 philosophers to be sitting simultaneously at the table.
  - Allow a philosopher to pick up the forks only if both are available (picking must be done in a critical section).
  - Use an asymmetric solution -- an odd-numbered philosopher picks up first the left chopstick and then the right chopstick. Even-numbered philosopher picks up first the right chopstick and then the left chopstick.