

# Operating Systems: Deadlocks

Neerja Mhaskar

Department of Computing and Software, McMaster University, Canada

**Acknowledgements:** Material based on the textbook Operating Systems Concepts (Chapters 5)

# System Model

- System consists of resources
- Resource types  $R_1, R_2, \dots, R_m$ 
  - *CPU cycles, memory space, I/O devices*
- Each resource type  $R_i$  has  $W_i$  instances.
  - Eg: You can have two instances of a printer available
- Each process utilizes a resource as follows:
  - **request**
  - **use**
  - **Release**
- if the resources not available, the process enters a *waiting state*.

# Deadlock

- **Deadlock** – when a waiting process has requested a resource held by other waiting processes.
- In a deadlock processes never finish executing
  - System resources are tied up.
  - Thus, preventing other processes from starting.


# Deadlock Characterization

Deadlock  $\Leftrightarrow$  below **four conditions hold** at the same time.

- **Mutual exclusion:** only one process at a time can use a resource.
- **Hold and wait:** a process holding at least one resource is **waiting to acquire additional resources** held by other processes.
- **No preemption:** a resource can be **released only voluntarily** by the process holding it, after that process finishes its task.

# Deadlock Characterization contd..

- **Circular wait:** there exists a set  $\{P_0, P_1, \dots, P_n\}$  of **waiting processes** such that  $P_0$  is waiting for a resource that is held by  $P_1$ ,  $P_1$  is waiting for a resource that is held by  $P_2$ , ...,  $P_{n-1}$  is waiting for a resource that is held by  $P_n$ , and  $P_n$  is waiting for a resource that is held by  $P_0$ .



circular-wait condition  $\Rightarrow$  the hold-and-wait condition.

# Resource-Allocation Graph

- **System resource-allocation graph** (directed graph

$G = (V, E)$ ) used to describe deadlocks.

- $V$  is partitioned into two sets:


- $P = \{P_1, P_2, \dots, P_n\}$ , the set consisting of all the processes in the system

- $R = \{R_1, R_2, \dots, R_m\}$ , the set consisting of all resource types in the system

- $E$  has two types of edges:


- **Request edge** – directed edge  $P_i \rightarrow R_j$ : Process  $P_i$  requested an instance of resource  $R_j$  and is waiting.

- **Assignment edge** – directed edge  $R_j \rightarrow P_i$ : Resource  $R_j$  assigned to process  $P_i$



# Resource-Allocation Graph (Cont.)


- Process




- Resource Type with 4 instances



- $P_i$  requests instance of  $R_j$




- $P_i$  is holding an instance of  $R_j$




# Resource Allocation Graph Examples

Resource allocation graph with a cycle but no Deadlock



Resource allocation graph with a cycle and a Deadlock



Cycles with circular wait condition satisfied

- $P_1 \rightarrow R_1 \rightarrow P_2 \rightarrow R_3 \rightarrow P_3 \rightarrow R_2 \rightarrow P_1$
- $P_2 \rightarrow R_3 \rightarrow P_3 \rightarrow R_2 \rightarrow P_2$



# Basic Facts

- If a graph contains no cycles  $\Rightarrow$  no deadlock
- If graph contains a cycle  $\Rightarrow$ 
  - if only one instance per resource type, then deadlock
  - if several instances per resource type, possibility of deadlock

# Methods for Handling Deadlocks

- Ensure that the system will **never** enter a deadlock state:
  - **Deadlock prevention:** ensure that atleast one of the necessary conditions (stated in slide #4) cannot hold.
    - Only the **Circular Wait condition is a practical** option for deadlock prevention.
  - **Deadlock avoidance:** requires that OS be given **additional information in advance** about the type and no. of resources a process will request. Based on this data OS **decides the waiting strategy for the process.**

# Methods for Handling Deadlocks

- **Deadlock detection:** Allow the system to enter a deadlock state and then recover.
- **IGNORE!** the problem and pretend that deadlocks never occur; used by most operating systems such as UNIX, Linux, and Windows!

# Deadlock Avoidance Strategy

- Requires that the OS be given **additional information in advance** about the type and no. of resources a process will request.
- Each process declares its ***maximum need = number of resources*** of each type that the process may need.
- The **deadlock-avoidance algorithm** dynamically examines the ***resource-allocation state*** to ensure that there can never be a **circular-wait condition**.

# Resource allocation state

**Resource-allocation state** is defined by

- Number of available resources
- Number of allocated resources
- the maximum need of all the processes in the system.

**Example:** Consider a system with a single resource type – magnetic tapes= 12.

Resource allocation state of the system at time  $T_0$ :

	<u>Maximum Need</u>	<u>Allocated resources</u>	<u>Available</u>
P0	10	5	12-9 = 3
P1	4	2	
P2	9	2	

# Safe State

- System is in **safe state** if there exists a sequence  $\langle P_1, P_2, \dots, P_n \rangle$  of ALL the processes in the system such that for each  $P_i$ :
  - Resources  $P_i$  needs  $\leq$  total available resources + resources held by/allocated to all  $P_1, P_2, \dots, P_{i-1}$ .
- We compute the needs of a process  $P_i$  as below
$$P_i\text{'s Need} = \text{Max. Need of Process } P_i - \text{Allocated resources for process } P_i$$

# Safe State Example

**Is the below system at time  $t_0$  in safe state?**

- Consider a system with 12 magnetic tapes.
- **Resource allocation state** of the system at time  $t_0$ :

	<u>Maximum Needs</u>	<u>Allocated resources</u>	<u>Available</u>
$P_0$	10	5	$12 - 9 = 3$
$P_1$	4	2	
$P_2$	9	2	

# Safe State Example – with Need vector

- $P_i$ 's Need = Max. Need of Process  $P_i$  – Allocated resources for process  $P_i$
- **Resource allocation state** of the system at time  $t_0$  with need vector:

	<u>Maximum Needs</u>	<u>Allocated resources</u>	<u>Need</u>	<u>Available</u>
$P_0$	10	5	5	3
$P_1$	4	2	2	
$P_2$	9	2	7	

Is there a sequence consisting of ALL processes ( $P_1, P_2, \dots, P_n$ ) in the system such that every process satisfies the safe state requirement?

*Safe state requirement is*

*Resources  $P_i$  needs  $\leq$  total available resources + resources held by/allocated to all  $P_1, P_2, \dots, P_{i-1}$*

YES! The sequence is  $\langle P_1, P_0, P_2 \rangle$



# Safe State Example Cont.

- Suppose at time  $t_1$  process  $P_2$  *requests an additional tape drive* and is assigned as well.
- **Resource allocation state** of the system at time  $t_1$ :


<u>Process</u>	<u>Maximum Needs</u>	<u>Allocated resources</u>	<u>Need</u>	<u>Available</u>
$P_0$	10	5	5	2
$P_1$	4	2	2	
$P_2$	9	3	6	

➤ Available resources at time  $t_1 = 2$

- Is the system at time  $t_1$  in safe state?

# Basic Facts

- If a system is in safe state  $\Rightarrow$  no deadlocks
- If a system is in unsafe state  $\Rightarrow$  possibility of deadlock
- Avoidance  $\Rightarrow$  ensure that a system will never enter an unsafe state.



# Deadlock Avoidance algorithm Outline

- Given the resource allocation state of a system,
  - The algorithm checks if the system is in a *safe state*.
  - If the system is in a safe state, whenever a process requests an instance of a resource type,
    - It checks to see if allocating the resources continues to have the system in a safe state.
      - If yes -- allocate the resources.
      - If no -- have the process wait.