

# Operating Systems: Deadlocks – PART II

Neerja Mhaskar

Department of Computing and Software, McMaster University, Canada

**Acknowledgements:** Material based on the textbook Operating Systems Concepts (Chapters 5)

# Deadlock Avoidance algorithm Outline


- Given the resource allocation state of a system,
  - The algorithm checks if the system is in a *safe state*.
  - If the system is in a safe state, whenever a process requests an instance of a resource type,
    - It checks to see if allocating the resources continues to have the system in a safe state.
      - If yes -- allocate the resources.
      - If no -- have the process wait.

# Deadlock Avoidance Algorithms

- Single instance of a resource type
  - Use a **resource-allocation graph**
- Multiple instances of a resource type
  - Use the **Banker's algorithm**


# Resource-Allocation Graph Scheme

- Add claim edges to existing resource allocation graph.
- **Claim edge**  $P_i \rightarrow R_j$  indicates that process  $P_i$  may request resource  $R_j$ ; *represented by a dashed line*
- Claim edge converts to request edge when a process requests a resource
- Request edge converted to an assignment edge when the resource is allocated to the process
- When a **resource is released by a process**, assignment edge reconverts to a claim edge
- Resources must be claimed *a priori* in the system



# Resource-Allocation Graph Algorithm

- Suppose that process  $P_i$  requests a resource  $R_j$
- The request can be granted **only if** converting the request edge to an assignment edge **does not result in the formation of a cycle** in the resource allocation graph
- We check for safety by using a cycle-detection algorithm.



## Unsafe State In Resource-Allocation Graph

Although  $R_2$  is free, we cannot allocate it to  $P_2$  since this will create a cycle!

# Banker's Algorithm Outline for *multiple instances* of a resource type

- The algorithm consists of two parts
  - **PART 1 - Safety Algorithm** – checks whether a system is in a safe state or not.
  - **PART 2 - Resource-Request Algorithm** – checks to see if resources requested by a process can be satisfied or not.
- Each process must a priori claim maximum use
- When a process gets all its resources it must return them in a *finite amount of time*
- When a process requests a resource, it *may have to wait*

# Data Structures for the Banker's Algorithm

Let  $n$  = number of processes, and  $m$  = number of resources types.

- **Available**: Vector of length  $m$ . If  $available[j] = k$ , there are  $k$  instances of resource type  $R_j$  available
- **Max**:  $n \times m$  matrix. If  $Max[i,j] = k$ , then process  $P_i$  may request at most  $k$  instances of resource type  $R_j$
- **Allocation**:  $n \times m$  matrix. If  $Allocation[i,j] = k$  then  $P_i$  is currently allocated  $k$  instances of  $R_j$
- **Need**:  $n \times m$  matrix. If  $Need[i,j] = k$ , then  $P_i$  may need  $k$  more instances of  $R_j$  to complete its task

$$Need[i,j] = Max[i,j] - Allocation[i,j]$$

# Example of Data Structures for Banker's Algorithm

- 5 processes  $P_0$  through  $P_4$ ;

3 resource types:

$A$  (10 instances),  $B$  (5 instances), and  $C$  (7 instances)

**Resource allocation state at time  $T_0$ :**

	Max. Need			Allocation			Available		
	A	B	C	A	B	C	A	B	C
$P_0$	7	5	3	0	1	0	3	3	2
$P_1$	3	2	2	2	0	0			
$P_2$	9	0	2	3	0	2			
$P_3$	2	2	2	2	1	1			
$P_4$	4	3	3	0	0	2			



# Need matrix for Banker's Algorithm

The content of the matrix **Need** is defined to be **Max. Need – Allocation**

- 5 processes  $P_0$  through  $P_4$ ;

3 resource types:

$A$  (10 instances),  $B$  (5 instances), and  $C$  (7 instances)

**Snapshot at time  $T_0$ :**

	Max. Need			Allocation			Need			Available		
	A	B	C	A	B	C	A	B	C	A	B	C
$P_0$	7	5	3	0	1	0	7	4	3	3	3	2
$P_1$	3	2	2	2	0	0	1	2	2			
$P_2$	9	0	2	3	0	2	6	0	0			
$P_3$	2	2	2	2	1	1	0	1	1			
$P_4$	4	3	3	0	0	2	4	3	1			

# Part 1 - Safety Algorithm outline explained with an example

**Step 1:** Maintain **Work** and **Finish** vectors of length  $m$  and  $n$ , respectively.

*Initialize:*

**Work** = **Available**

**Finish** [ $i$ ] = **false**,  $i \in [0, n-1]$

Work. = (3 3 2)

Finish = 

False	False	False	False	False
-------	-------	-------	-------	-------

**Step 2:** Find a process  $P_i$  such that both:

(a) **Finish** [ $i$ ] = **false**

(b) **Need** <sub>$i$</sub>  ≤ **Work**

$P_1$  satisfies conditions (a) and (b)

If no such  $i$  exists, go to step 4

## Part 1 - Safety Algorithm outline explained with an example – contd...

**Step 3:**

**$Work = Work + Allocation_i$**

**$Finish[i] = true$**

**Go to step 2**

$Work = Work + Allocation_1$

$= (3\ 3\ 2) + (2\ 0\ 0) = (5\ 3\ 2)$

$Finish[1] = true$

$Finish =$ 

False	True	False	False	False
-------	------	-------	-------	-------

Next, we see that process  $P_3$ ,  $P_4$ ,  $P_2$ , and  $P_0$  all satisfy the conditions in step 2.

**Step 4:** If  **$Finish[i] == true$**  for all  $i$ , then the system is in a safe state

$Finish =$ 

True	True	True	True	True
------	------	------	------	------

Therefore, the system is in a safe state and the sequence of processes satisfying the safety requirement is -

**$\langle P_1, P_3, P_4, P_2, P_0 \rangle$**

## Part – 2 Resource-Request Algorithm for Process $P_i$

- **$Request_i$**  = request vector for process  $P_i$ . If  **$Request_i[j] = k$**  then process  $P_i$  wants  $k$  instances of resource type  $R_j$
- **Step 1:** If  **$Request_i \leq Need_i$**  go to step 2. Otherwise, raise error condition, since process has exceeded its maximum claim
- **Step 2:** If  **$Request_i \leq Available$** , go to step 3. Otherwise,  $P_i$  must wait, since resources are not available
- **Step 3:** Pretend to allocate requested resources to  $P_i$  and update the system state as follows:

$$Available = Available - Request_i;$$

$$Allocation_i = Allocation_i + Request_i;$$

$$Need_i = Need_i - Request_i;$$

- If safe  $\Rightarrow$  the resources are allocated to  $P_i$
- If unsafe  $\Rightarrow P_i$  must wait, and the old resource-allocation state is restored

## PART 2 - Resource-Request Algorithm Explained with an Example

- **$P_1$  requests resources (1 0 2)**
- Check if  **$\text{Request}_1 \leq \text{Need}_1$** 
  - $(1\ 0\ 2) \leq (1\ 2\ 2) \Rightarrow \text{true}$
- Check if  **$\text{Request}_1 \leq \text{Available}$** 
  - $(1\ 0\ 2) \leq (3\ 2\ 2) \Rightarrow \text{true}$
- Pretend that resources requested have been granted.
- Update system state. Max need,  $\text{Allocation}_1$  and  $\text{Need}_1$  data structures
  - $\text{Available}_1 = (3\ 2\ 2) - (1\ 0\ 2) = (2\ 2\ 0)$
  - $\text{Allocation}_1 = (2\ 0\ 0) + (1\ 0\ 2) = (3\ 0\ 2)$
  - $\text{Need}_1 = (1\ 2\ 2) - (1\ 0\ 2) = (0\ 2\ 0)$

## PART 2 - Resource-Request Algorithm Explained with an Example contd...

**Updated resource allocation state:**

	Max. Need			Allocation			Need			Available		
	A	B	C	A	B	C	A	B	C	A	B	C
P <sub>0</sub>	7	5	3	0	1	0	7	4	3	2	2	0
P <sub>1</sub>	3	2	2	3	0	2	0	2	0			
P <sub>2</sub>	9	0	2	3	0	2	6	0	0			
P <sub>3</sub>	2	2	2	2	1	1	0	1	1			
P <sub>4</sub>	4	3	3	0	0	2	4	3	1			

## PART 2 - Resource-Request Algorithm Explained with an Example contd...

- Run safety algorithm on the **updated resource allocation state**.
- System is in safe state and the sequence of processes satisfying the safety requirement is  **$\langle P_1, P_3, P_4, P_2, P_0 \rangle$**

### Updated Resource allocation state after request has been granted for $P_1$

	Max. Need			Allocation			Need			Available		
	A	B	C	A	B	C	A	B	C	A	B	C
$P_0$	7	5	3	0	1	0	7	4	3	2	2	0
$P_1$	3	2	2	3	0	2	0	2	0			
$P_2$	9	0	2	3	0	2	6	0	0			
$P_3$	2	2	2	2	1	1	0	1	1			
$P_4$	4	3	3	0	0	2	4	3	1			

## Resource-Request Algorithm Example – Cont...

- When system in this state, can request for (3 3 0) by  $P_4$  be granted?
  - Check if **Request<sub>4</sub> ≤ Available**
    - $(3\ 3\ 0) \leq (2\ 2\ 0) \Rightarrow \text{false}$
    - The request cannot be granted.
- When system in this state, can request for (0 2 0) by  $P_0$  be granted?
  - Check if **Request<sub>0</sub> ≤ Need<sub>0</sub>**
    - $(0\ 2\ 0) \leq (7\ 4\ 3) \Rightarrow \text{true}$
  - Check if **Request<sub>0</sub> ≤ Available**
    - $(0\ 2\ 0) \leq (2\ 2\ 0) \Rightarrow \text{true}$
  - Pretend to grant the resources requested.



## Resource-Request Algorithm Example – Cont...

### Updated Resource allocation state

	Max. Need			Allocation			Need			Available		
	A	B	C	A	B	C	A	B	C	A	B	C
P <sub>0</sub>	7	5	3	0	3	0	7	2	3	2	0	0
P <sub>1</sub>	3	2	2	3	0	2	0	2	0			
P <sub>2</sub>	9	0	2	3	0	2	6	0	0			
P <sub>3</sub>	2	2	2	2	1	1	0	1	1			
P <sub>4</sub>	4	3	3	0	0	2	4	3	1			


- However, since no sequence of processes exist that satisfies the safe state requirement, P<sub>0</sub>'s request cannot be granted as doing so will leave the system in an unsafe state.

# Deadlock Detection


- Allow system to enter deadlock state
- Use deadlock detection algorithm to check if a deadlock exists
- If deadlock exists, use a recovery scheme to recover from the deadlock

# Deadlock Detection - Single Instance of Each Resource Type

- A variant of the resource-allocation graph if all resources have only **a single instance** - used for deadlock detection
- Nodes are processes, and an edge  $P_i \rightarrow P_j$  in the wait for graph implies that  $P_i$  is waiting for  $P_j$  to release a resource that  $P_j$  needs.




Resource-Allocation Graph




Corresponding wait-for graph

# Deadlock Detection - Single Instance of Each Resource Type

***If a cycle exists in the wait-for graph, then the system is in deadlock.***



Resource-Allocation Graph



Corresponding wait-for graph

# Deadlock Detection Algorithm for Multiple Instances of a Resource Type

- The algorithms needs to know
  - The number of *available resources and allocated resources* for each resource type.
  - The number of *requested resources* by all processes in the system.
- Given the above,
  - The deadlock detection algorithm checks whether the system is ***in a deadlocked*** state or not.
  - If the system is in a deadlocked state, then the algorithm also identifies the processes involved in the deadlock.

# Deadlock Detection Algorithm for Multiple Instances of a Resource Type Data Structures

- **Available:** A vector of length  $m$  indicates the number of available resources of each type
- **Allocation:** An  $n \times m$  matrix defines the number of resources of each type currently allocated to each process
- **Request:** An  $n \times m$  matrix indicates the current request of each process. If  $Request[i][j] = k$ , then process  $P_i$  is requesting  $k$  instances of resource type  $R_j$ .

# Deadlock Detection Algorithm

1. Let **Work** and **Finish** be vectors of length **m** and **n**, respectively

Initialize:

(a) **Work** = **Available**

(b) For  $i = 1, 2, \dots, n$ ,

i. if **Allocation<sub>i</sub> ≠ 0**, then **Finish[i] = false**;

ii. otherwise, **Finish[i] = true**

2. Find an index **i** such that both:

(a) **Finish[i] == false**

(b) **Request<sub>i</sub> ≤ Work**

If no such **i** exists, go to step 4

*Items in red highlight the differences in Deadlock detection algorithm and the safety algorithm described under Banker's algorithm.*

# Detection Algorithm (Cont.)

3.  **$Work = Work + Allocation_i$**

**$Finish[i] = true$**

go to step 2

4. If  **$Finish[i] == false$** , for some  $i$ ,  $1 \leq i \leq n$ , then the system is in deadlock state. Moreover, if  **$Finish[i] == false$** , then  **$P_i$**  is deadlocked.



# Example for Deadlock Detection Algorithm

- Five processes  $P_0$  through  $P_4$
- Three resource types  
A (7 instances), B (2 instances), and C (6 instances)
- Snapshot of the system at time  $T_0$ :

	Allocation			Request			Available		
	A	B	C	A	B	C	A	B	C
$P_0$	0	1	0	0	0	0	0	0	0
$P_1$	2	0	0	2	0	2			
$P_2$	3	0	3	0	0	0			
$P_3$	2	1	1	1	0	0			
$P_4$	0	0	2	0	0	2			

# Example of Detection Algorithm Cont...

## Step 1:

1.  $Work = Available = (0\ 0\ 0)$

Since  $Allocation_i \neq 0$ , for all  $i \in [1, n]$

Finish =	False	False	False	False	False
----------	-------	-------	-------	-------	-------

## Step 2: Find an index $i$ such that both:

1.  $Finish[i] == false$
2.  $Request_i \leq Work$

If no such  $i$  exists, go to step 4

**$P_0$  satisfies the above two conditions.**

# Example of Detection Algorithm Cont...

- Step 3:

- $\text{Work} = \text{Work} + \text{Allocation}_0 = (0\ 0\ 0) + (0\ 1\ 0) = (0\ 1\ 0)$

- $\text{Finish}[1] = \text{true}$

- go to step 2

- We see that process  $P_2$ ,  $P_3$ ,  $P_1$ , and  $P_4$  all satisfy the conditions in step 2.

- Finally, in **Step 4: Finish** = 

True	True	True	True	True
------	------	------	------	------

- Therefore, the system is ***not in a deadlocked state***, as the following sequence of processes results in all values of the **Finish** vector to be **True**:

- $\langle P_0, P_2, P_3, P_1, P_4 \rangle$

# Example of Detection Algorithm Cont...

- Suppose  $P_2$  requests an additional instance of type C.
- Then below is the updated snapshot of the system including this request:

	Allocation			Request			Available		
	A	B	C	A	B	C	A	B	C
$P_0$	0	1	0	0	0	0	0	0	0
$P_1$	2	0	0	2	0	2			
$P_2$	3	0	3	0	0	1			
$P_3$	2	1	1	1	0	0			
$P_4$	0	0	2	0	0	2			

- Can reclaim resources held by process  $P_0$ , but insufficient resources to fulfill other processes' requests
- **Deadlock exists, consisting of processes  $P_1$ ,  $P_2$ ,  $P_3$ , and  $P_4$**

# Recovery from Deadlock

## ■ Process Termination

- Abort all deadlocked processes
- Abort one process at a time until the deadlock cycle is eliminated

## ■ Resource Preemption

- **Selecting a victim** – minimize cost
- **Rollback** – return to some safe state, restart process for that state
- **Starvation** – same process may always be picked as victim. Possible solution - include number of rollback in cost factor