

Operating Systems: Main Memory – Part II

Neerja Mhaskar

Department of Computing and Software, McMaster University, Canada

Acknowledgements: Material based on the textbook Operating Systems Concepts (Chapter 9)

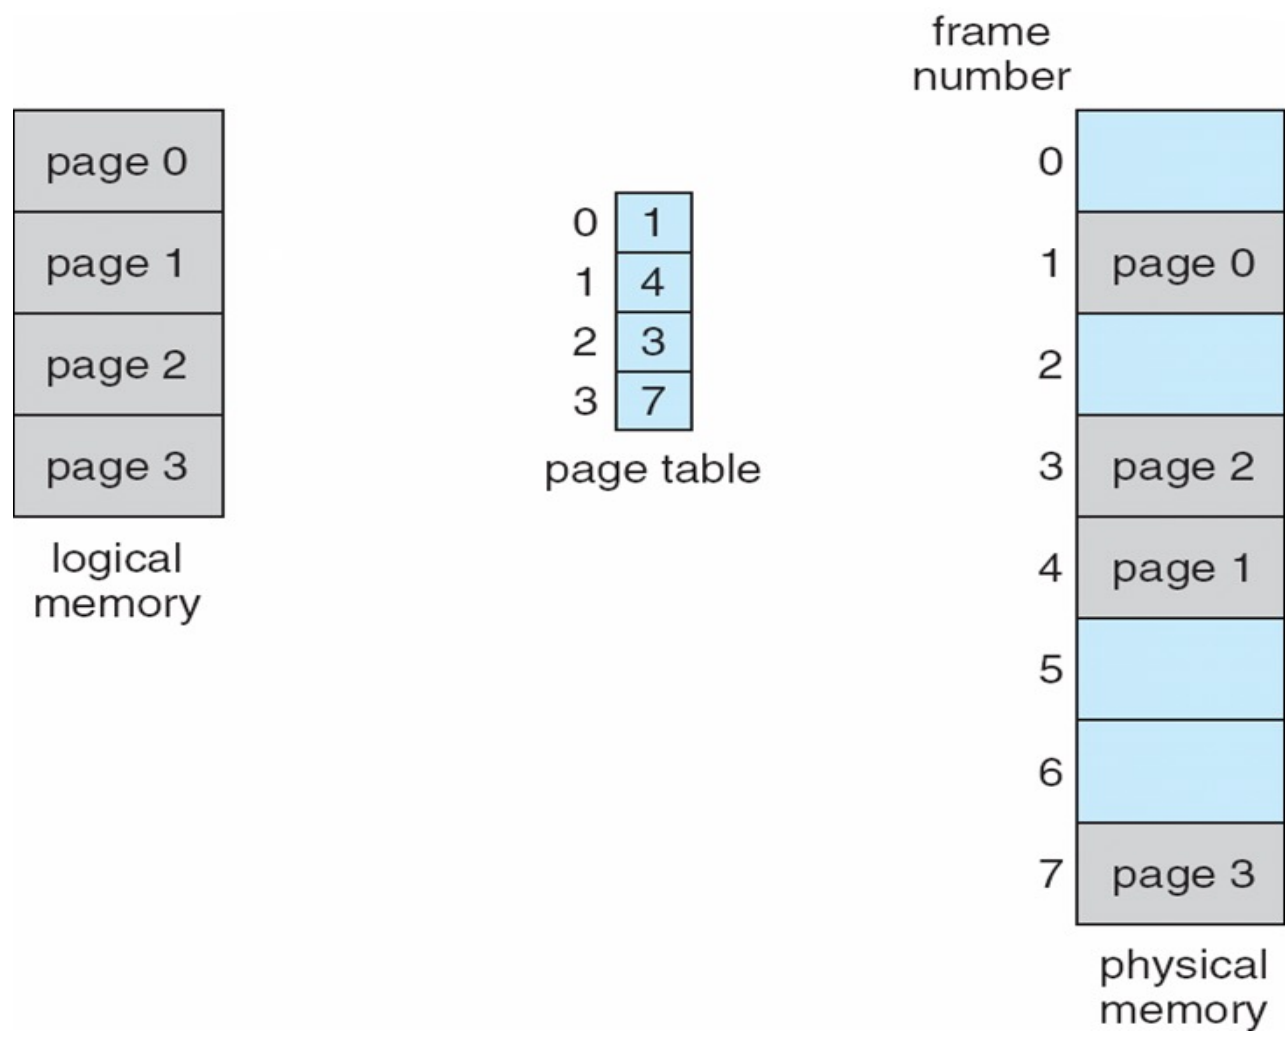
Paging

- **Paging** - a memory-management scheme that permits a process's physical address space to be **non-contiguous**.
- Logical address space of a process is divided into fixed blocks called **pages**
- Physical address; that is main memory is divided into fixed-sized blocks called **frames**.
- **Page size = Frame size**

Paging

- Pages of a process are scattered in main memory and stored in frames.
- Each executing process has its own **page table**.
- **Page table** – keeps track of the pages residing in main memory
 - Indexed by page numbers
 - Stores the frame number in which the page loaded.

Paging Model of Logical and Physical Memory



Paging – Some facts

- Page size are generally 4 KB or 8 KB (may range from 4KB – 1GB).
 - Power of 2 as it makes address translation easy.
- The page size (like the frame size) is defined by the hardware

Advantages:

- Simple to implement.
- Avoids external fragmentation but suffers from internal fragmentation.

Logical address and paging

- Logical address space = $2^4 = 16$ bytes.
- What is the range of addresses in this logical space?
- How many bits do you need to represent a logical address in this space?

Logical address and paging


Logical address space = $2^4 = 16$ bytes.

- Range of addresses in this logical space = $0..15(2^4-1)$
- Bits needed to represent an address in this space = 4
- If the page size is be 4 bytes, then the logical address space is divided into 4 pages.

Logical addresses	Logical address in binary
0	0000
1	0001
2	0010
3	0011
4	0100
5	0101
6	0110
7	0111
8	1000
9	1001
10	1010
11	1011
12	1100
13	1101
14	1110
15	1111

Paging - Address Translation

- Logical address now consists of two parts:



- **Page number (p)** – used as an index into a **page table**
 - **Page offset (d)** – is the offset/displacement within the page.
- Page offset - ranges from 0 to page size -1
 - Page size = frame size => page offset = frame offset.
 - **Physical memory address = frame number + page offset.**

Address Translation Scheme - Generalized

For given logical address space size = 2^m , page size = 2^n and physical address space = 2^r

- Number of bits to represent logical address =
- Number of bits to represent physical address =
- Number of bits to represent offset =
- Number of pages =
- Number of bits to represent a page number =
- Number of frames =
- Number of bits to represent a frame number =

page number	page offset
p	d
m - n	n

Paging and Address Translation Example

Logical address space = $2^4 = 16$ bytes

- Number of bits to represent logical address =
- Page size = $2^2 = 4$ bytes
- Number of bits to represent offset =
- Number of bits to represent a page number =
- Total # of pages =

Memory/Physical address size = $2^5 = 32$ bytes

- Number of bits to represent physical address =
- Total # of frames =
- Number of bits to represent a frame number =

Logical addresses	P number	d (offset)
0	00	00
1	00	01
2	00	10
3	00	11
4	01	00
5	01	01
6	01	10
7	01	11
8	10	00
9	10	01
10	10	10
11	10	11
12	11	00
13	11	01
14	11	10
15	11	11

Paging and Address Translation Whole Picture!

- CPU generates Logical address **3 [0011]** – page# - 00 and offset = 11
- Page 0 resides in the frame# 5 [101]
- Logical address 3 maps to physical address = **10111**(5*4 + 3) = 20 + 3 = 23.

0	a
1	b
2	c
3	d
4	e
5	f
6	g
7	h
8	i
9	j
10	k
11	l
12	m
13	n
14	o
15	p

logical memory

0	5
1	6
2	1
3	2

page table


0	
4	i j k l
8	m n o p
12	
16	
20	a b c d
24	e f g h
28	

physical memory

Frames in Memory


- The operating system manages physical memory.
 - Needs to know which frames are allocated, which frames are available, how many total frames there are, etc.
 - To answer the above questions OS maintains a data structure called a **frame table (free-frame list)**.

Allocating Frames



(a)

Before allocation



(b)

After allocation

Implementation of Page Table

- For small size page tables – provide dedicated registers.
- For large page tables – **page table kept in main memory**
 - **Page-table base register (PTBR)** points to the starting address of the page table in main memory.
 - In this scheme every data/instruction access requires two memory accesses
 - One for the page table and one for the data / instruction!

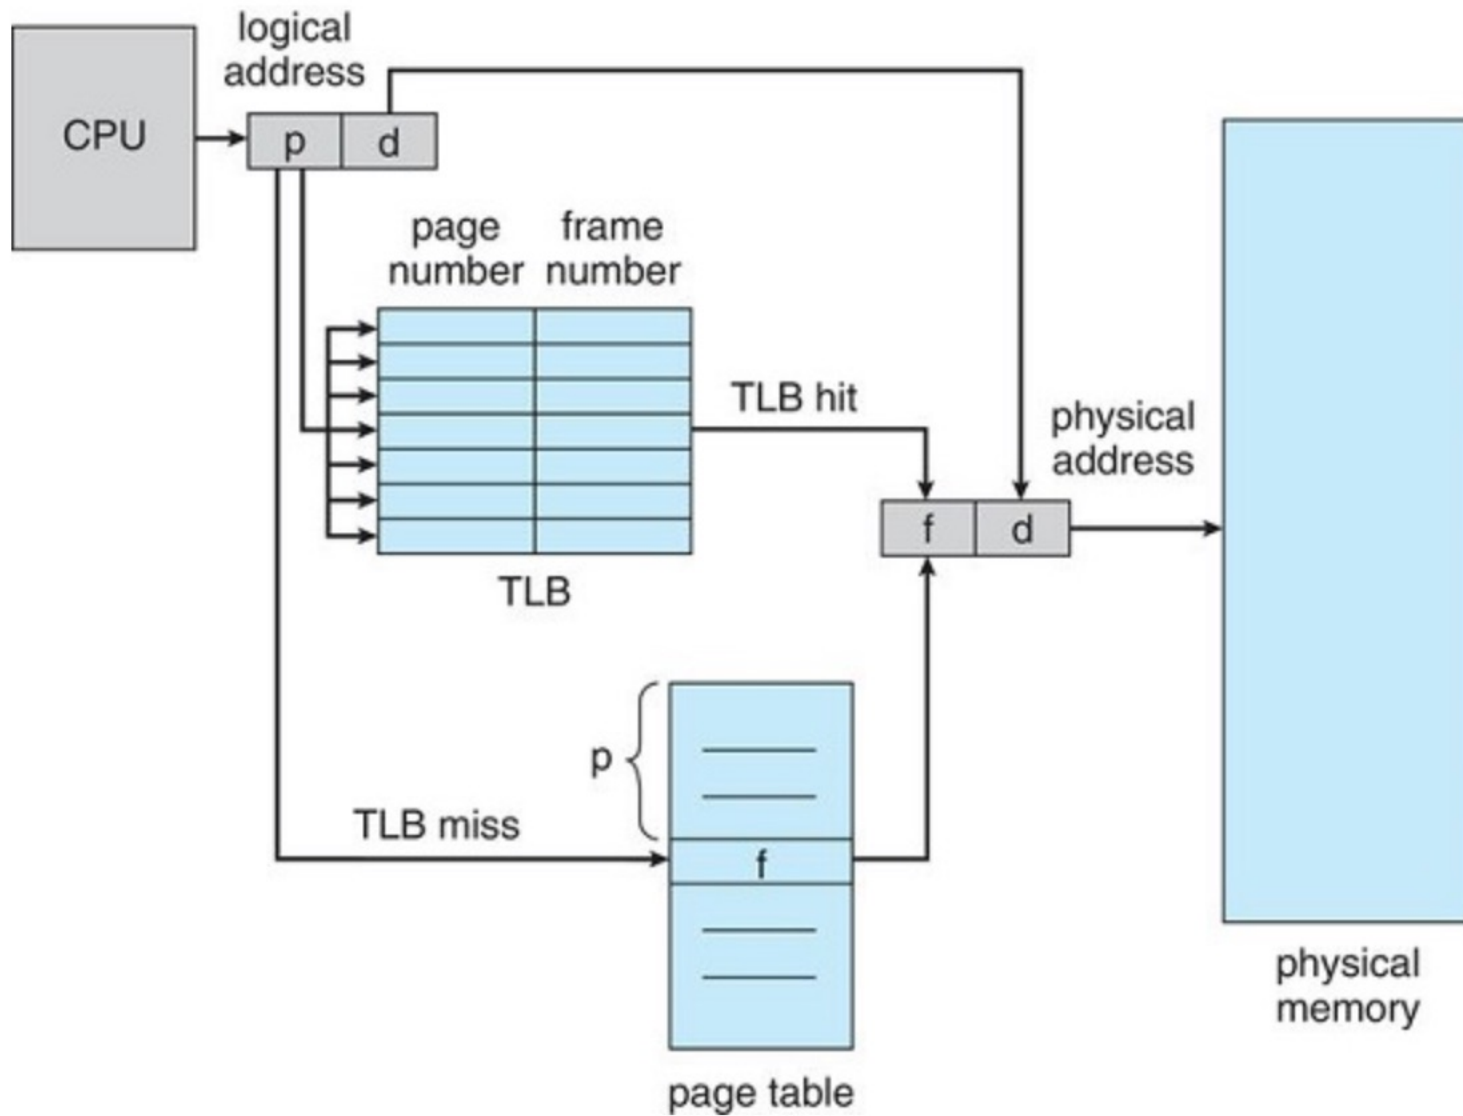
Implementation of Page Table

- To mitigate the two-memory access problem current machines, use a special fast-lookup hardware cache called **translation look-aside buffers (TLBs)**
 - TLB is associative, high-speed memory.
 - TLBs typically are small (64 to 1,024 entries)

Translation Look-aside Buffer (TLB)

- When a logical address is generated by the CPU, its page number is presented as a data item to TLB.
- This data is checked with all the entries in TLB **simultaneously**.
- If page is found (called a **TLB Hit**) its corresponding frame# is returned.
- If page is not found in TLB (called a **TLB Miss**) page table is searched for the page and its corresponding frame# is returned and entry for the page added in TLB.
- Some TLBs store **address-space identifiers (ASIDs)** in each TLB entry – uniquely identifies each process to provide address-space protection for that process
 - Otherwise need to flush at every context switch

Paging Hardware With TLB



TLB - Effective Access Time

- **Hit ratio (α):** is the percentage of times a page number of interest is found in TLB.
- Memory access time = m
- TLB access time = t
- $EAT = (m + t) \alpha + (2m + t)(1 - \alpha)$

TLB - Effective Access Time Example

- **Suppose memory access time (m) = 100ns**
- If we fail to find the page in TLB then accessing memory is $2m = 200\text{ns}$.
- Ignoring TLB access time t
 - If hit-ratio (α) = 80%
 - $\text{EAT} = 0.80 \times 100 + 0.20 \times 200 = 120\text{ns}$
 - Consider more realistic hit ratio $\rightarrow \alpha = 99\%$
 - $\text{EAT} = 0.99 \times 100 + 0.01 \times 200 = 101\text{ns!}$

Structure of the Page Table

- Consider a 32-bit logical address space as on modern computers, and page size of 4 KB (2^{12})

Page number | Page offset

20 bits	12 bits
----------------	----------------

- No of entries in a single level page table = 2^{20} entries
- If each page entry is 4 bytes:
 - Size of page table = $4 * 2^{20} = 4 \text{ MB!}$
 - Don't want to allocate that contiguously in main memory!

Structure of the Page Table Cont...

- Common techniques for structuring a page table:

- **Hierarchical paging**

- Hashed page table (used for logical address spaces $> 2^{32}$)

- Inverted page table

Multi- Level/Hierarchical Paging

- **Multi- Level/Hierarchical Paging** – is a technique in which the page table is paged; that is, broken down into multiple page tables of the same size as the page.
- Need an outer page table to track these pages!
- With varying sizes of the logical address space and pages, we could have a very large outer page table.
 - In this case, we repeat this process by paging the outer page table (thus introducing another layer of page table).

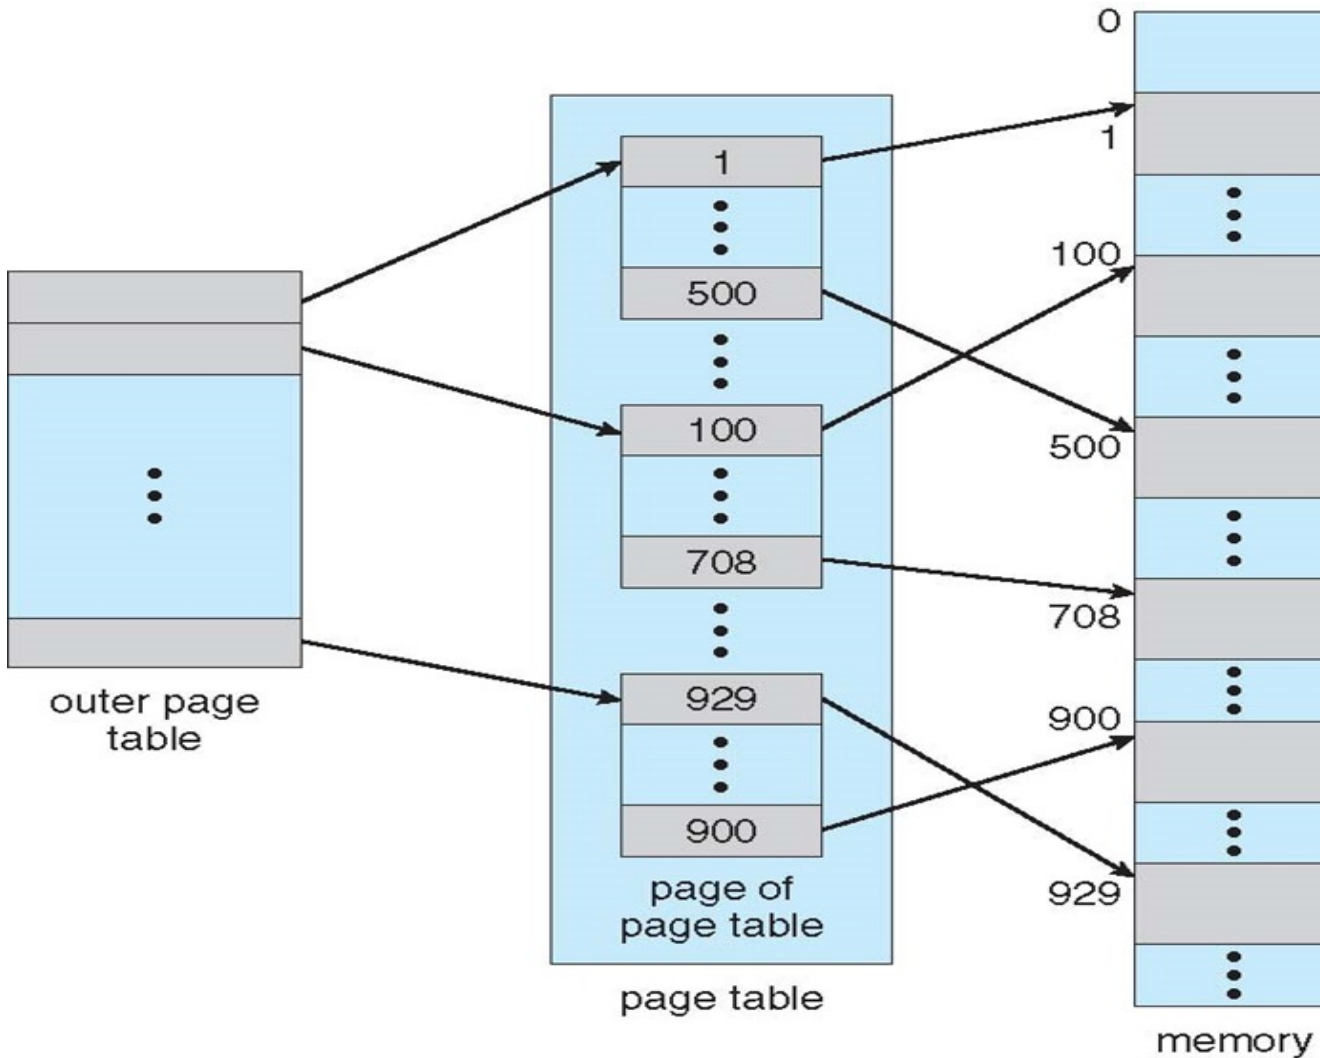
Multi- Level/Hierarchical Paging Example

- From the example previously shown the size of a single level page table = $4 * 2^{20}$ bytes.
- Number of pages of the page table =
Size of the single level page table/size of a page
= $4 * 2^{20} / 2^{12} = 2^{10}$

Multi- Level/Hierarchical Paging

- To track these pages of the page table we need a **second level page table** (or outer page table)!
- Size of this outer page table
 - = Number of pages of the page table * page entry size
 - = $2^{10}(2^2) = 2^{12}$ bytes = size of one page.
- Since there are two levels of page tables, this above technique is called **two-level paging**.

Two-Level Page-Table Scheme Illustration



Address Translation in two level paging scheme - Example

On a 32-bit machine with 4K page size (previous example)


No. of page table entries in a page is (at most) = $2^{12}/4 = 2^{10}$.

The distribution of no. of bits to represent each level of the multi-level page table, and the page offset is as follows:


10	10	12
----	----	----

Address Translation in two level paging scheme - Generalized

- Since the page table is paged, the page number is further divided into:
 - p_1 is the displacement/offset in the outer page table, and
 - p_2 is the displacement/offset within the page of the inner page table
- Thus, a logical address is as follows:



Address Translation in two level paging scheme



- The starting address of the outer page table is stored in **PTBR**.
- p_1 is the offset into this outer page table.
- We extract the address at this memory address, which is the starting address of the page of the page table.
- Then, p_2 is the offset into this inner page to read the frame number.
- Once we get the frame number, we use the offset d to compute the address of the actual instruction/data.

Hierarchical paging example

- Logical address space = 2^5 bytes
- Page size = 2^2 bytes
- Physical address space = 2^8 bytes
- Page table entry size = 8 bits = 1 bytes
- No. of pages to represent this logical space = 2^3
- Size of a single level page table = $2^3 * 1 = 2^3$
- No. of entries of a page of a page table = $2^2/1 = 2^2$
- No. of pages required to represent the page table = $2^3/2^2 = 2$.
- 1 outer page table needed to hold the addresses of these to pages.

Hierarchical paging example continued..

■ Address Translation Scheme:

- 1 bit to represent offset in the outer page table
- 2 bit to represent offset in the inner page table
- 2 bit to represent offset within the page.

1	2	2
---	---	---

Multi Level Hierarchical Paging

- Consider a system with logical address space 64-bit logical address space with page size 4KB (2^{12})
- If the page table entry size is 4 bytes, then number of page table entries in a page = $2^{12}/4 = 2^{10}$
- We need 10 bits to represent an offset in this page of a page table.
- Therefore, according to the two-level paging scheme the address translation scheme is as follows:

42	10	12
-----------	-----------	-----------

- The size of the outer page table = $2^{42} * 4 = 2^{44}$ bytes --- too big to store in a single outer page of the page table!

Multi Level Hierarchical Paging

- Two-level paging scheme is no longer appropriate!
- Three level page scheme; that is, page the outer page table?
 - Any page of the page table (inner or outer) can have only 2^{10} entries. Therefore, we will need 10 bits to represent this offset.

32	10	10	12
----	----	----	----

- 2nd outer page also too big!
- Keep paging the different levels of outer page table, till it can be contained in a single page.

2	10	10	10	10	10	12
---	----	----	----	----	----	----

Inverted Page Table

- Rather than each process having a page table and keeping track of all possible logical pages, track all **physical pages (that is frames)**
- **Inverted page table** has one entry for each **frame** of memory
- Entry consists of the virtual address of the page stored in that real memory location, with information about the process that owns that page.

Example: Consider a computer system with a 32-bit logical address, 4-KB page size, and 24-bit physical address.

- A conventional, single-level page table contains 2^{20} entries.
- An inverted page table contains **2^{12} entries!**