

Operating Systems: File System Interface, Implementation, and System Internals – Part I

Neerja Mhaskar

Department of Computing and Software, McMaster University, Canada

Acknowledgements: Material based on the textbook Operating Systems Concepts (Chapter 13, 14 and 15)

File Systems

- **File System** - provides efficient and convenient access to disk by allowing data to be stored, located and retrieved easily.
- The file system consists of two distinct parts:
 - A collection of files
 - A directory structure
- Information on disk is stored in the form of **files** and organized using **directories**.
- Each file system has its own logic and data structures that reside on **secondary storage (disks)** to store and organize files.

File System Examples

- File system examples

- UFS (Unix file system) - UNIX

- Ext3/Ext4 (extended file system) - LINUX

- FAT, NTFS etc. – Windows

File

- **File** is the smallest logical storage unit in a file system and is a collection of related information defined by its creator.
 - A file is a sequence of bits, bytes, or a more complex data item.
 - It consists of **logical records** that range from a byte to several bytes.
 - Typically file contents are broken down into **physical block sizes** (usually 512 bytes or more)
 - File systems I/O performed in group of blocks called **clusters**.

Internal File Structure

- Disks have sectors that determine the physical block size (**which is the same or a multiple of the sector size**).
- All disk I/O is performed in units of these physical blocks. which are stored on one or more sectors.
- Physical block is not always equal to the **logical record of a file**.
 - Solution: Packing a number of logical records into physical blocks!

Internal File Structure

- In UNIX all files are streams of bytes.
 - Hence the logical record size is 1 byte.
 - The file system automatically packs and unpacks bytes into physical blocks—say, 512 bytes per block—as necessary.

Internal File Structure

- The packing of logical records can be done either by the user's application program or by the operating system.
- Because disk space is always allocated in blocks, all file systems suffer from **internal fragmentation**

File

- A file is associated with the following:
- **File attributes** – Name, Identifier (e.g., inode number), text etc.
- **File operations** – Create, write, read, delete, truncate, repositioning within a file.
- **File type** – Executable (.bin, .exe), text file (.txt, .doc)

File Control Block

- Every file has a **file control block (FCB)** that contains all the information about the file (e.g.: ownership, permissions, and location of the file contents.)
- FCB in Unix/Linux systems is called an **inode**.

file permissions
file dates (create, access, write)
file owner, group, ACL
file size
file data blocks or pointers to file data blocks

A typical file-control block

File structure

- Different file types have different file structures.
 - These file structures are used by Operating systems and applications programs to work with them.
 - **All operating systems must support the executable file structure to load and run programs.**
- Supporting different file structure can increase the size of the OS making it hard to manage and work with.
 - Therefore, Unix treats all files as sequence of bytes!

File Access Methods

- File access methods

- **Sequential Access:** Information in the file is processed in order, one record after the other.
 - Most common
 - Example: editors and compilers are accessed sequentially.
- **Direct access:** Information in the file is accessed directly (and in no particular order).
 - Example: in databases data is accessed directly.

Directories

- **Directories** – enable segregating files into groups and managing them as groups.
- The directory contains information about the files, including attributes, location, and ownership.
- **Unix treats directories and files as same entities**
- Operations that are to be performed on a directory:
 - Operations regarding a file: Create, delete rename, search.
 - Traverse the file system and list directory contents.
- Various schemes for defining logical structure of a directory exist.
 - Single-level, two-level, tree-structure, Acyclic-Graph Directories and General Graph directories.

Protection

- Files must be kept safe for reliability (against accidental damage), and protection (against deliberate malicious access.)
- Protection is achieved by controlling access to a file.
- A file can be accessed in various ways:
 - **Read, Write, Execute, Append, and Delete**

Access control lists (ACL)

- Is a protection mechanism in which each file/directory has a list of users and the type of access permitted for each user.
- Main issue with ACL is the length.
 - Solution: Condense the list by classifying users.

Access Lists and Groups in Linux

- **Classes of users on Unix / Linux:** Owner, Group, and Others.
- **Mode of access:** R (read), W (write), and X (execute)
- UNIX uses a set of 9 access control bits, in three groups of three
- Each class of user is given RWX permission using 3 bits (bit value 0 = deny access and 1 to grant access).
- File/Directory permissions set using `chmod` command.
- Example: Suppose we want to set the below permissions on a file named `test.txt`:

		RWX
a) owner access	7 \Rightarrow	1 1 1
b) group access	6 \Rightarrow	1 1 0
c) public access	1 \Rightarrow	0 0 1

- The command in UNIX is \Rightarrow `chmod 761 test.txt`

Question


What access rights does the below command specify on the file test.txt?

```
chmod 550 test.txt
```

Disk Structure – Partitions and Volumes

- Disks can be broken up into multiple **partitions, slices, or mini-disks**,
 - Each partition is called a logical disk which can have its own filesystem.
- Multiple physical disks can be combined into one **volume**, i.e., a larger virtual disk
 - This volume can have its own filesystem spanning the physical disks.
- Each volume containing file system also tracks that file system's info in **device directory** or **volume table of contents**.

A Typical File-system Organization



File System Mounting

- Before your computer can use any kind of storage device (such as a hard drive, CD-ROM, or network share), it must be made accessible through the computer's file system.
- This process is called **mounting**.
- An (unmounted) file system is mounted at a **mount point** - the location (**typically, an empty directory**) within the computer's file system where the file system is to be attached.

File-System Implementation

- Several on-disk and in-memory data structures are used to implement a file system

On Disk Structures

- **Boot control block (boot block in Unix)** is per volume
 - contains information needed by the system to boot OS from that volume.
Needed only if volume contains OS and is usually first block of volume.
- **Volume control block (superblock in UNIX)** is per volume
 - Contains volume (or partition) details (e.g.: Total # of blocks, # of free blocks, block size, free block pointers or array etc.)

File-System Implementation

On Disk Structures continued:

- **Directory structure** is per file system
 - Used to organizes the files. In Unix file system, it contains file names and associated inode numbers.
- **File Control Block, FCB,** (per file) containing details about the file.

File-System Implementation - In-Memory File System Structures

In-memory structures are used for both file-system management and performance improvement via caching.

- **Mount table** - contains information about each mounted volume
- **In-memory directory-structure cache** - holds the directory information of recently accessed directories.

File-System Implementation - In-Memory File System Structures

- **System-wide open-file table** contains a copy of the FCB of each open file, as well as tracks the number of processes that have the file open.
- **Per-process open-file table** contains a pointer to the appropriate entry (for the file) in the system-wide open-file table and some other fields.
- **Buffers** hold file-system blocks when they are being read from disk or written to disk.

Creating a new file

- An application program calls the file system to create a file.
- The file system knows the format of the directory structures.
- To create a new file, it **allocates a new FCB**.
- The system then **reads the appropriate directory into memory**.
 - Updates it with the new file name and FCB identifier, and writes it back to the disk.


Opening/Reading a file

- The `open()` call passes a file name to the logical file system.
- The file system first searches the system-wide open-file table to check if the file is already in use by another process.
 - YES => create per-process open-file table entry
 - NO => Search directory structure for the given file name. Once the file is found, its FCB is copied into the system-wide open-file table in memory.
 - System-wide open-file table stores the FCB, and also keeps tracks of the number of processes that have the file open.

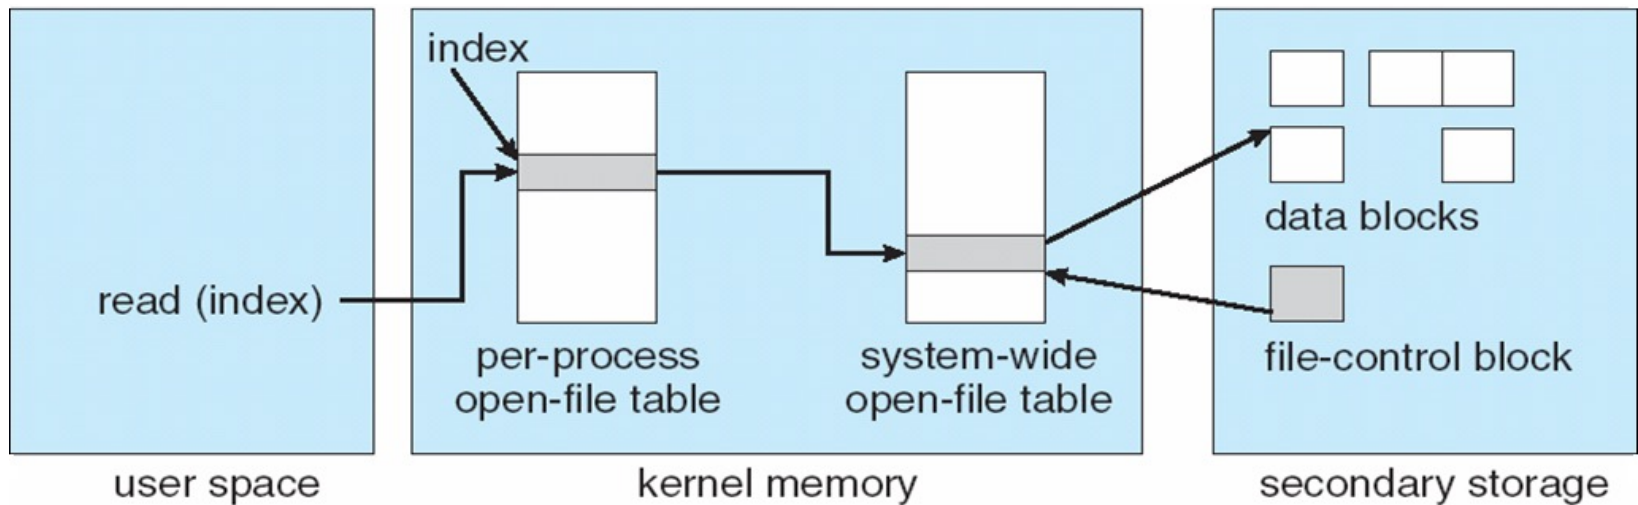
Opening/Reading a file

- **Update per-process open-file table** to store the pointer to the entry in the system-wide open-file table and some other fields (current location, access mode).
- returns a pointer to the appropriate entry (matching the filename) in the per-process file-system table. All file operations use this pointer (UNIX calls it **file descriptor**).

Opening a file – In Memory File System Structures



Reading a file – In Memory File System Structures



Closing a file

- When a process closes the file, the per-process table entry is removed
- The system-wide entry's open count is decremented.