

**Question 1****0 / 1 point**

Consider a dataset with  $p$  features and  $n$  observations. Which of the following statements is true for Principal Component Analysis?

- a ☐ There are at most  $\max(n + 1, p)$  principal components.
- b ☐ There are at least  $\min(n - 1, p)$  principal components.
- c ☐ There are at most  $\max(n - 1, p)$  principal components.
- d ☐ There are at most  $\min(n - 1, p)$  principal components.

**Question 2****0 / 1 point**

Which of the following hyperparameters controls the size of the region of the input image that the convolutional filter processes at each step?

- a ☐ Batch size
- b ☐ Number of epochs
- c ☐ Kernel size
- d ☐ Learning rate

**Question 3****0 / 1 point**

Which of the following is a major limitation of the K-Means clustering algorithm?

- a ☐ K-Means cannot handle non-numeric data.
- b ☐ K-Means can only be used for clustering spherical clusters.
- c ☐ K-Means requires the user to specify the number of clusters after the model is trained.
- d ☐ K-Means is sensitive to the initial placement of centroids.

**Question 4****0 / 1 point**

What is the primary goal of Principal Component Analysis (PCA)?

- a ☐ To increase the number of dimensions in the dataset to improve performance.
- b ☐ To reduce the dimensionality of the data while preserving as much variance as possible.
- c ☐ To remove outliers from the dataset.
- d ☐ To classify the data into different categories based on features.

**Question 5****0 / 1 point**

Which of the following types of problems is Naive Bayes best suited for?

- a ☐ Problems where classes are not linearly separable.
- b ☐ Problems where features are strongly correlated.
- c ☐ Text classification tasks such as spam email detection.
- d ☐ Tasks involving a large number of continuous features with no categorical variables.

## Question 6

0 / 1 point

Which of the following statements best describes the primary difference between **generative** and **discriminative** algorithms?

- a ☐ Discriminative models can create new data instances, while generative models focus on modeling different kinds of instances.
- b ☐ Discriminative models require fewer features than generative models to achieve the same level of accuracy.
- c ☐ Discriminative models learn the probability of the classes given the data, while generative models learn the joint probability of the data and the class.
- d ☐ Generative models typically perform better on large datasets, while discriminative models are best for smaller datasets.

### Question 7

0 / 1 point

For what tasks in Machine Learning is Visualization useful?

- a ☐ Understanding training data
- b ☐ Inspecting the model
- c ☐ Communicating model results
- d ☐ Dealing with high dimensional data
- e ☐ All of the above

### Question 8

0 / 1 point

Which of the following statements is **incorrect** for Principal Component Analysis (PCA)?

- a ☐ PCA is used for feature selection.
- b ☐ Principal Component directions are computed using the order sequence of eigenvectors.
- c ☐ For multi-dimensional data, there are multiple principal components.
- d ☐ PCA considers linear combination of all features.

**Question 9****0 / 1 point**

Which of the following is **NOT** a core component of a Markov Decision Process (MDP)?

- a ☐ States (S)
- b ☐ Loss Function (L)
- c ☐ Rewards (R)
- d ☐ Actions (A)

**Question 10****0 / 1 point**

What is the key feature of a Recurrent Neural Network (RNN) that distinguishes it from traditional feedforward neural networks?

- a ☐ RNNs process data sequentially, maintaining a memory of previous inputs through feedback connections.
- b ☐ RNNs have no internal state and do not capture temporal dependencies.
- c ☐ RNNs use a fixed-size architecture with no connections between layers.
- d ☐ RNNs are only used for image classification tasks.

**Question 11****0 / 1 point**

Which of the following is a key assumption made by the Naive Bayes algorithm?

- a ☐ Features are conditionally independent given the class label.
- c ☐ Features are highly correlated with each other.
- d ☐ The class distribution follows a Gaussian distribution.
- e ☐ The dataset contains only binary features.

**Question 12****0 / 1 point**

Which of the following is a variant of Recurrent Neural Networks (RNNs) designed to address the vanishing gradient problem?

- a ☐ Long Short-Term Memory (LSTM)
- b ☐ Convolutional Neural Network (CNN)
- c ☐ Generative Adversarial Network (GAN)
- d ☐ Autoencoder



**Question 13****0 / 1 point**

What does the Kullback-Leibler (KL) divergence measure?

- a ☐ The similarity between two probability distributions.
- b ☐ The difference between the mean values of two distributions.
- c ☐ The distance between two data points in Euclidean space.
- d ☐ The variance within a single distribution.

**Question 14****0 / 1 point**

What is a policy in the context of Reinforcement Learning?

- a ☐ A measure of how much reward the agent will receive in the future.
- b ☐ A function that determines the agent's reward based on its actions.
- c ☐ A set of rules that govern how an agent should interact with the environment.
- d ☐ A strategy that defines the probability of taking a specific action given a state.

## Question 15

0 / 1 point

What does the Bellman Equation represent in the context of Reinforcement Learning?

- a ☐ A recursive relationship that helps calculate the optimal state-value functions for the agent.
- b ☐ A method to calculate the transition probabilities between states.
- c ☐ A function used to update the agent's position in the environment.
- d ☐ A formula that calculates the reward for each action taken by an agent.

## Question 16

0 / 1 point

Consider a 5x5 input image with the following values:

1	2	3	4	5
6	7	8	9	10
11	12	13	14	15
16	17	18	19	20
21	22	23	24	25

Consider you are applying a 3x3 filter (kernel) with the following values:

1	0	-1
1	0	-1
1	0	-1

Perform a **valid** convolution operation, which means the filter is applied only where it fits entirely inside the input image. What is the resulting value at the **top-left** corner of the output feature map?

- a ☐ -12
- b ☐ 6
- c ☐ -6
- d ☐ 12



## Question 17

0 / 1 point

Consider the below equation for reinforcement learning, where  $R$  is the reward at state  $t$  ( $T$  is the terminal state) and there are  $k$  states to sweep over. With  $G$  being the overall reward during learning, what is gamma ( $\gamma$ )?

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots + \gamma^{k-1} R_{T+k+1}$$

- a ☐ Number of Actions
- b ☐ Discount Rate
- c ☐ Learning Rate
- d ☐ Iteration Number

**Question 18****0 / 1 point**

Which of the following is a key characteristic of hierarchical clustering?

- a ☐ It uses a predefined number of clusters to begin with and iterates to optimize them.
- b ☐ It requires the user to specify the number of clusters before starting the algorithm.
- c ☐ It starts with all data points as individual clusters and merges them iteratively.
- d ☐ It starts by treating all data points as one single cluster and splits them iteratively.

**Question 19****0 / 1 point**

Which of the following layers is typically used in Convolutional Neural Networks to reduce the spatial dimensions of feature maps, while keeping the most important information?

- a ☐ Convolutional Layer
- b ☐ Fully Connected Layer
- c ☐ ReLU Layer
- d ☐ Pooling Layer

**Question 20****0 / 1 point**

In Reinforcement Learning, what does an agent aim to maximize?

- a ☐ The cumulative reward over time.
- b ☐ The total number of actions taken.
- c ☐ The total number of states available.
- d ☐ The number of states visited.

## Answers

1) d

2) c

3) d

4) b

5) c

6) c

7) e

8) a

9) b

10) a

11) a

12) a

13) a

14) d

15) a

16) c

17) b

18) c

19) d

20) a